

Exercise Session 3

Systems Programming and Computer Architecture

Fall Semester 2022

Agenda

- Last week's assignment
- More on C-Programing
 - File I/O

Last Week's Assignment


Bitwise Operations

Bit Lab – bitCount(x)


- Naïve Approach:
 - $(x \& 1) + ((x \gg 1) \& 1) + \dots + ((x \gg 31) \& 1)$
 - requires: 31 +, 32 &, 31 >> \Rightarrow 94 operators!!
- Another Idea:
 - Divide 32 bits into segments of «bit counters»




accumumulate bit 25 to 32



accumumulate bit 17 to 24



accumumulate bit 9 to 16



accumumulate bit 1 to 8

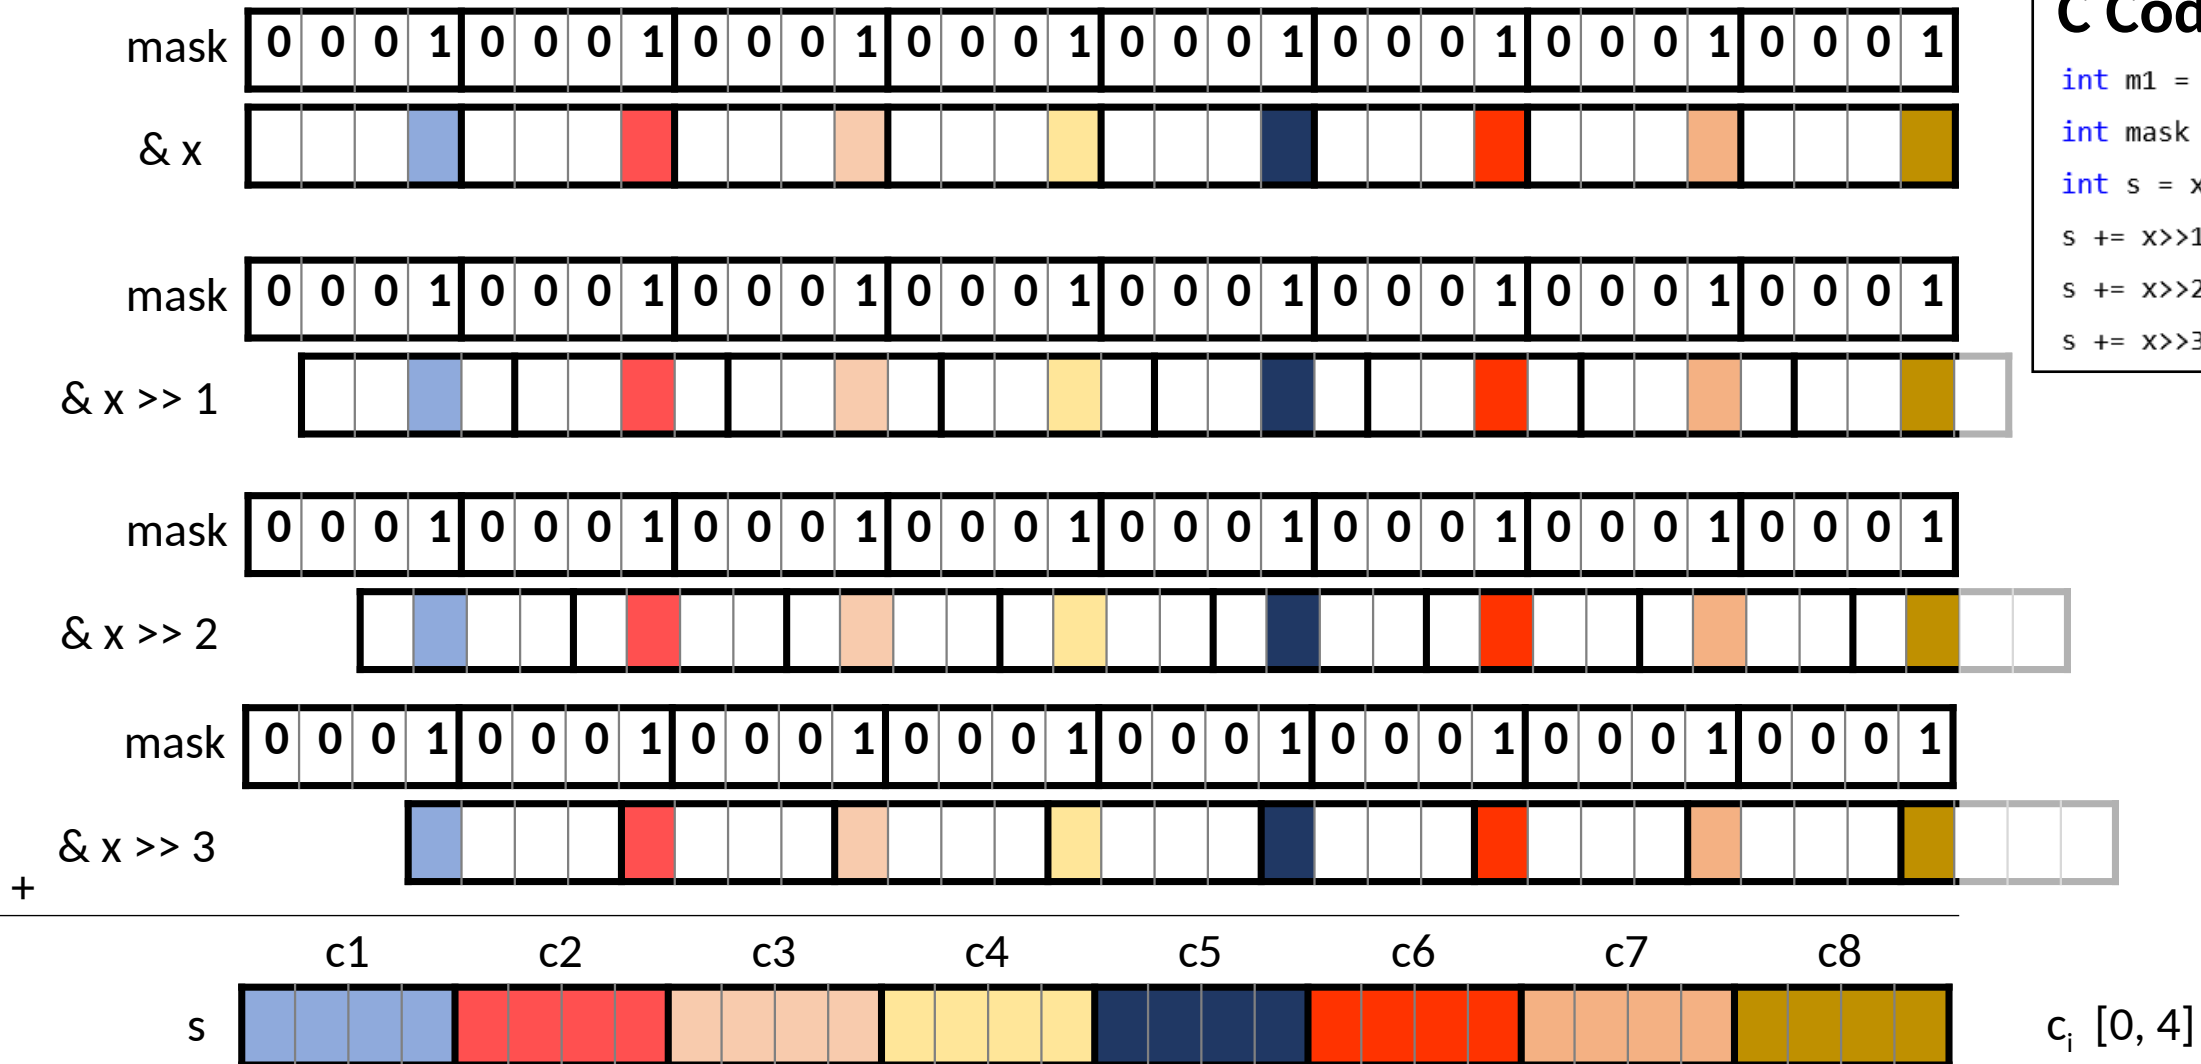
Bit Lab - bitCount(x) with counters

- Mask: `unsigned int m = 1 + (1 << 8) + (1 << 16) + (1 << 24)` 6

0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---
- Accumulate:
`unsigned int counts = (x & m) + ((x >> 1) & m) + ... ((x >> 7) & m)` 22
- Sum up:
`unsigned int sum = (counts & 0xFF) + ((counts >> 8) & 0xFF) +
((counts >> 16) & 0xFF) + ((counts >> 24) & 0xFF)` 10
- # Operators? 38

Bit Lab - bitCount(x) – better solution?

Accumulate in eight 4-bit Counters (rather than 4 8-bit Counters)




C Code

```
int m1 = 0x11 | (0x11 << 8);  
int mask = m1 | (m1 << 16);  
int s = x & mask;  
s += x >> 1 & mask;  
s += x >> 2 & mask;  
s += x >> 3 & mask;
```

Step 1: Reduce from 8 to 4 Counters

Combine 4 upper counters with 4 lower counters




Goal 1 Counter: $c_1 + c_2 + c_3 + c_4 + c_5 + c_6$

C Code

```
s = s + (s >> 16);
```

Step 2: Reduce from 4 to 2 Counters

Combine 1st with 2nd and 3rd with 4th



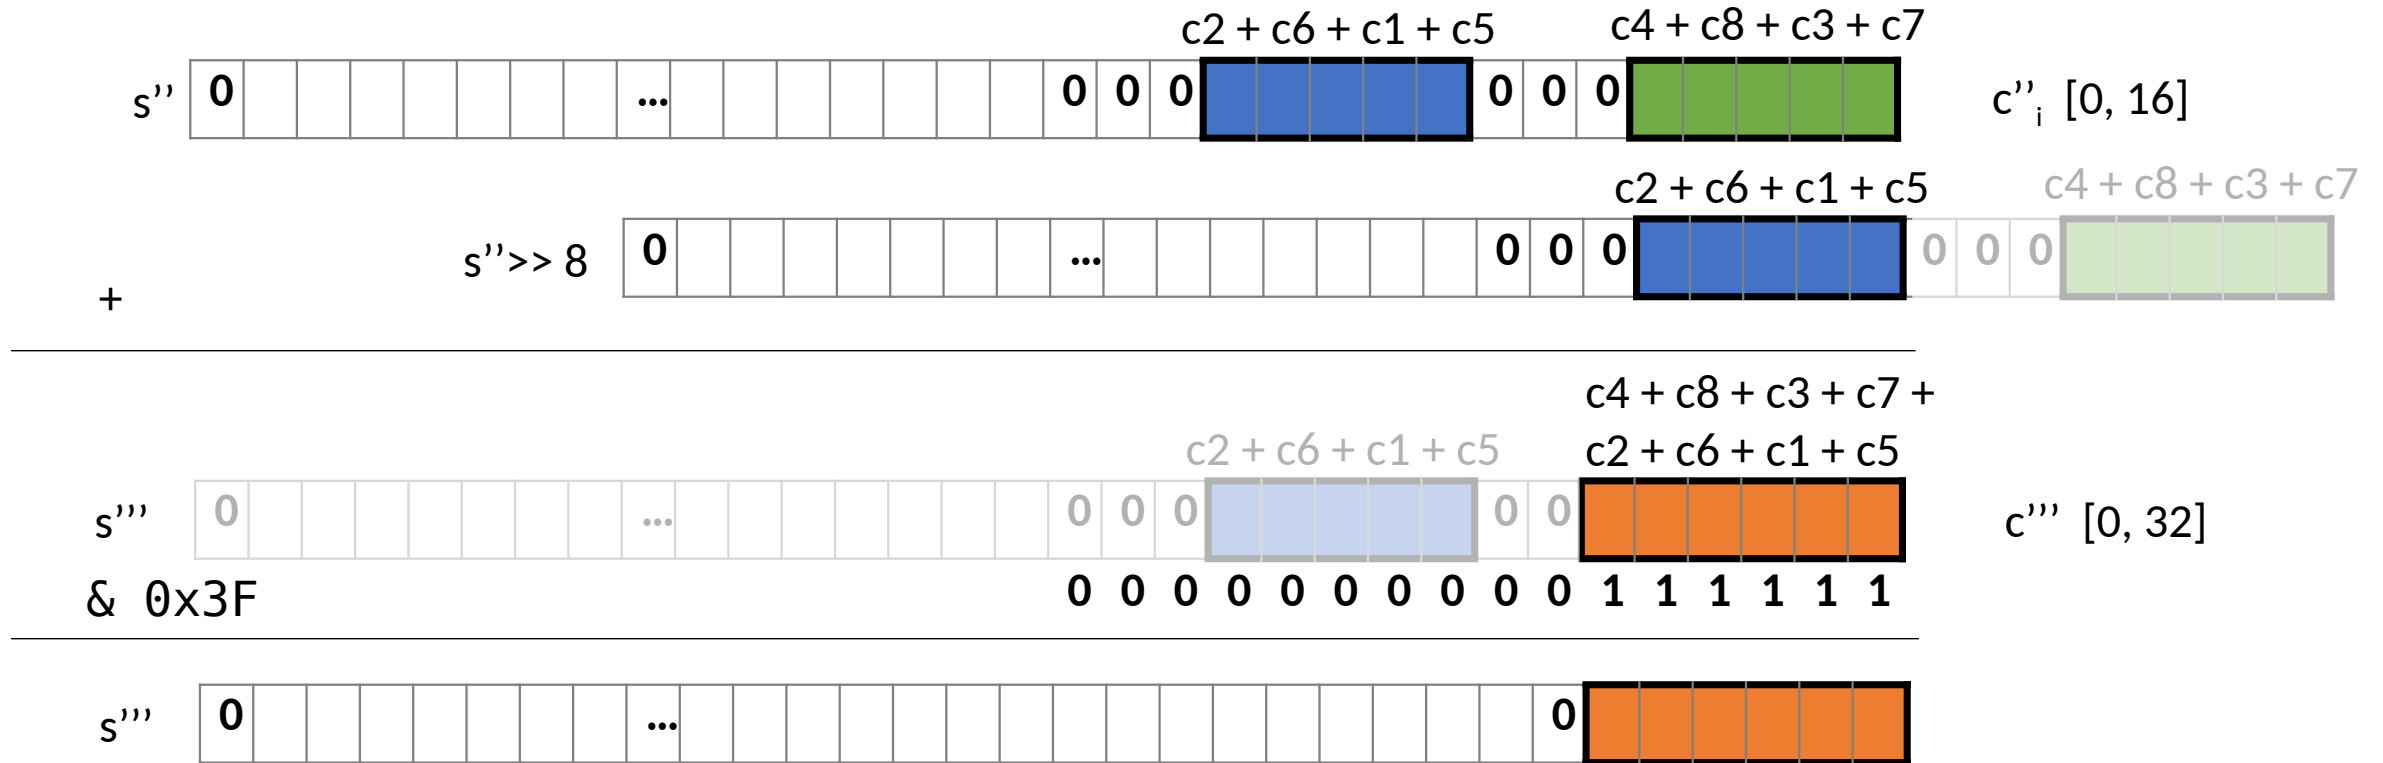
Goal 1 Counter: $c_1 + c_2 + c_3 + c_4 + c_5 + c_6$

C Code

```
mask = 0xF | (0xF << 8);  
s = (s & mask) + ((s >> 4) & mask);
```


Step 3: Reduce from 2 to 1 Counter

Combine the two remaining counters





Goal 1 Counter: $c1 + c2 + c3 + c4 + c5 + c6$

C Code

```
return (s + (s>>8)) & 0x3F;
```

Bit Lab - bitCount(x) – better solution?

```
int bitCount(int x) {  
    /* Sum 8 groups of 4 bits each */  
    int m1 = 0x11 | (0x11 << 8);  
    int mask = m1 | (m1 << 16);  
    int s = x & mask;  
    s += x>>1 & mask;  
    s += x>>2 & mask;  
    s += x>>3 & mask;  
    /* Now combine high and low order sums */  
    s = s + (s >> 16);  
    /* Low order 16 bits now consists of 4 sums,  
       each ranging between 0 and 8.  
       Split into two groups and sum */  
    mask = 0xF | (0xF << 8);  
    s = (s & mask) + ((s >> 4) & mask);  
    return (s + (s>>8)) & 0x3F;  
}
```



Operators: 25

Bit Lab

Questions for Assignment 1?

More on C-Programming

and short overview of
assignment 02

Me:

I am good in C language.

Interviewer:

Then write "Hello World" using C.

Me:



HELLO
WORLD

Integer Data Types

- You may use the types defined in C99's `stdint.h`

```
#include <stdint.h>

void main(int argc, char *argv[]) {
    uint8_t x = 22;
    int16_t y = 4434;
    int64_t z = 1 << 44;
}
```

- Keep in mind not to overflow by using a too small representation for the value

<http://www.cplusplus.com/reference/cstdint/>

<http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/stdint.h.html>

Reverse an array

Write a C program that has a function that:

- accepts an array of 32-bit unsigned integers and a length
- reverses the elements of the array in place
- returns void (nothing)

```
#include <stdint.h>
...
void reverse(uint32_t *arr, unsigned int len) {
    ...
}
```

Choose Your Types Wisely

```
void reverse_array(int array[], unsigned int length)
{
    for (int i = 0; i < length; i++) {
        /* ... */
    }
}
```

Use the -Wall flag every time !!!!

```
test.cc: In function 'void reverse_array(int,
unsigned)':
test.cc:6:22: warning: comparison between signed and
unsigned integer expressions [-Wsign-compare]
    for (int i = 0; i < length; i++) {
```

Box-and-arrow diagram

- Use a box-and-arrow diagram for the given program to explain what it prints out.
- Pen & Paper Exercise: hand it in manually or scan it

ETH zürich

Pointer arithmetic

```
#include <stdio.h>

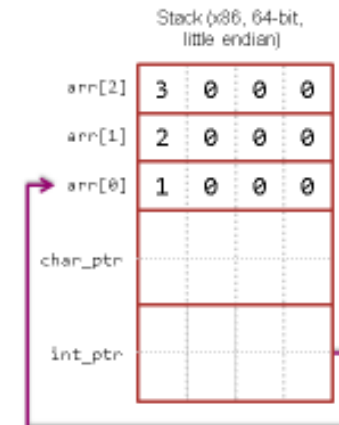
int main(int argc, char **argv) {
    int arr[3] = {1, 2, 3};
    int *int_ptr = &arr[0];
    char *char_ptr = (char *) int_ptr; // more "type casting" later

    printf("int_ptr: %p; *int_ptr: %d\n", int_ptr, *int_ptr);
    int_ptr += 1;
    printf("int_ptr: %p; *int_ptr: %d\n", int_ptr, *int_ptr);
    int_ptr -= 2; // uh oh
    printf("int_ptr: %p; *int_ptr: %d\n", int_ptr, *int_ptr);

    printf("char_ptr: %p; *char_ptr: %d\n", char_ptr, *char_ptr);
    char_ptr += 1;
    printf("char_ptr: %p; *char_ptr: %d\n", char_ptr, *char_ptr);
    char_ptr += 2;
    printf("char_ptr: %p; *char_ptr: %d\n", char_ptr, *char_ptr);

    return 0;
}
```

pointerarithmetic.c



See lecture slides (04 – Pointers)!

Little vs. big endian

Write a C program that prints out whether the computer it is running on is little endian or big endian. (hint: pointers and casts)

```
// returns true if little endian, 0 if big endian
bool is_little_endian() {
    ...
}
```

Initializing Memory

(First whitespace-separated word)

It's not safe to assume malloc() returns zeroed memory. So initialize!

```
#include <string.h>
```

```
int *data = (int*)malloc(10 * (sizeof(int)));  
if (data == NULL) {  
    printf("No memory");  
    exit(1);  
}  
// ensure memory is zeroed out  
memset(data, 0, 10 * sizeof(int));
```

Remember, use compiler flags

```
gcc <FILES> -Wall -Wpedantic -Wextra -Werror -std=c99 -Wmissing-  
prototypes
```



Complex numbers

For this you have to implement a complex number module with following functions:

- add
- subtract
- multiply
- divide

As well as a test driver, which contains the main function.

Dynamic memory allocation

Implement the two functions `alloc_set()` and `free_set()`, but be careful during allocations:

```
struct person * alloc_person() {  
    struct person * ret = (struct person *)malloc(sizeof(struct person));  
    ret->first_name = (char*)malloc(100 * sizeof(char));  
    ret->last_name = (char*)malloc(100 * sizeof(char));  
    return ret;  
}
```

- `malloc()` can fail, returns `NULL`.
- Correct code must not leak memory.

Cleanup code

```
struct person * alloc_person() {  
    struct person * ret = (struct person *)malloc(sizeof(struct person));  
    if(!ret)  
        return NULL;  
    ret->first_name = (char*)malloc(100 * sizeof(char));  
    if(!ret->first_name) {  
        free(ret);  
        return NULL;  
    }  
    ret->last_name = (char*)malloc(100 * sizeof(char));  
    if(!ret->last_name) {  
        free(ret->first_name);  
        free(ret);  
        return NULL;  
    }  
    return ret;  
}
```

WC

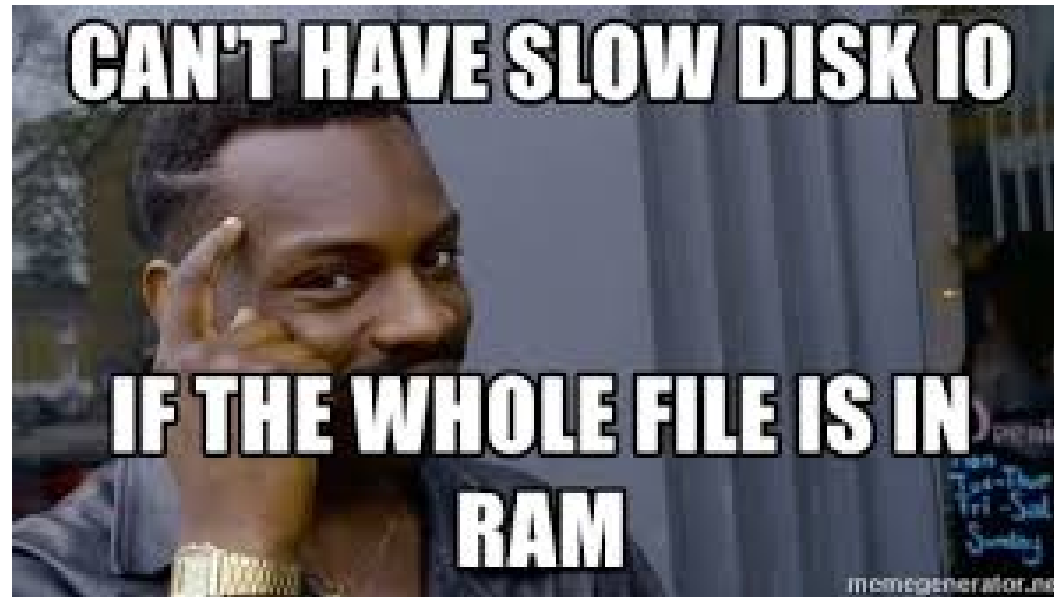
What is wc?

wc is a Unix utility that displays the count of characters, words and lines present in a file.

Implement this unix utility step by step while solving one problem at a time.

Start from the given shell in wc.c and then just add the missing components.

File I/O (File descriptors, read, write)

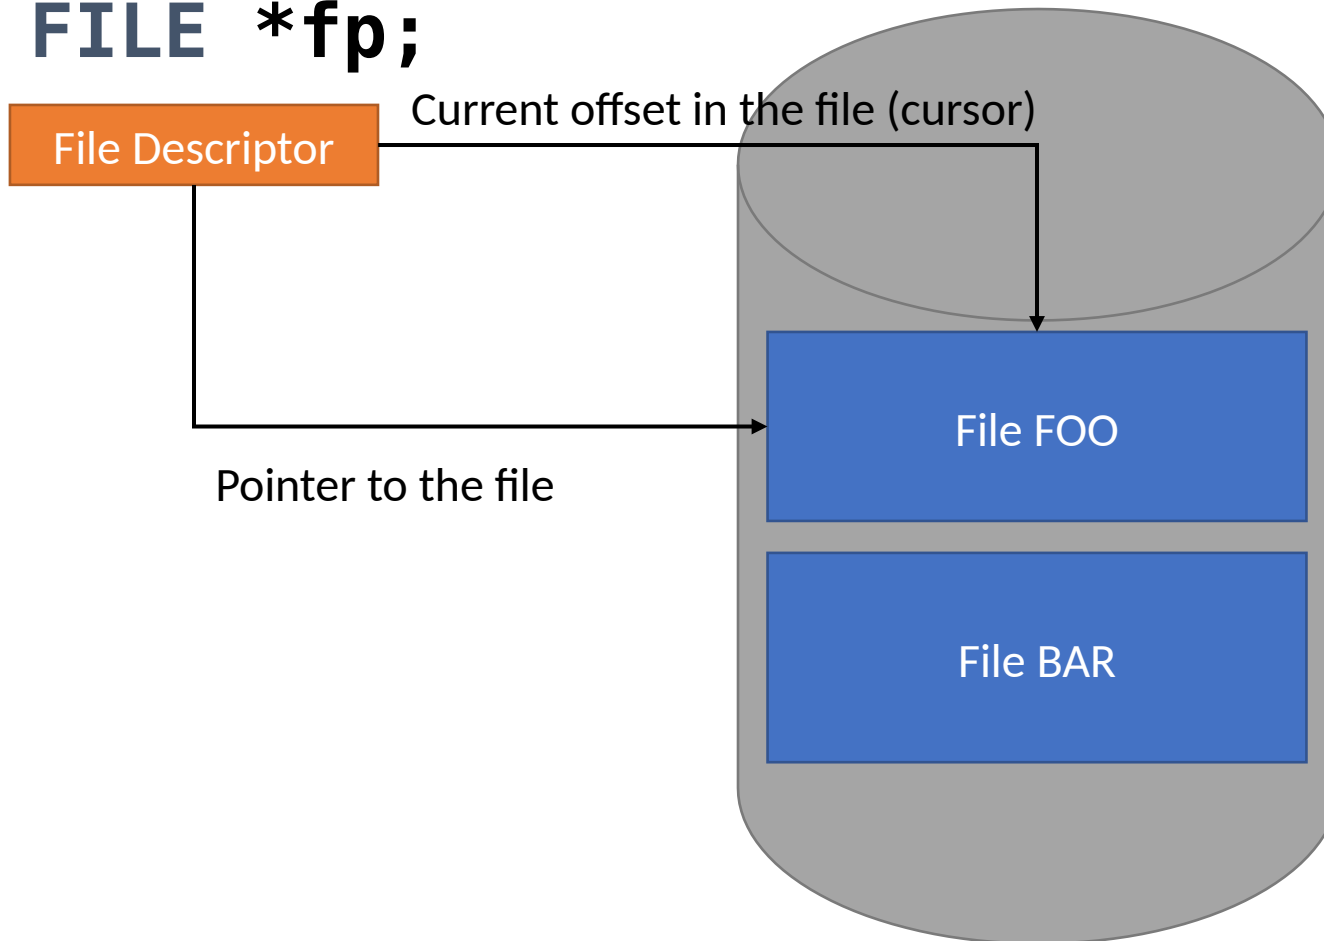


Accessing Files

- In general you cannot access the file directly
- You need support from the operating system to open/read/write/close a file.
 - This is called a system call (syscall)
 - > Lecture Computer Systems
- All file related declarations are in the `stdio.h`
`#include <stdio.h>`

The File Descriptor

FILE *fp;



- To get a file descriptor you have to open a file

```
FILE *fopen(const char  
*filename,  
const char *mode);
```

- The opening may fail. Check return value!
- Close the file in the end

```
int fclose(FILE *fp);
```

(Text) File Opening Modes

Mode	Read	Write	File Not Exists	File Exists
r	Yes, from beginning	No	Error	FILE* descriptor returned
w	No	Yes, from the beginning	New file created, FILE* Descriptor returned	FILE* descriptor returned
a	No	Yes, from the end (append)	New file created, FILE* descriptor returned	FILE* descriptor returned
r+	Yes, from beginning	Yes, from beginning	Error	FILE* descriptor returned.
w+	Yes, from beginning	Yes, From beginning	New file created, FILE* descriptor returned	Delete file contents (overwrite), File* descriptor returned
a+	Yes, from beginning	Yes, from the end (append only)	New file created, FILE* descriptor returned	FILE* descriptor returned
.. b	Binary flag that can be added for binary IO			

(Text) File Opening Modes

Mode	Read	Write	File Not Exists	File Exists
r	Yes, from beginning		Error	FILE* descriptor returned
w	No		New file created, FILE* descriptor returned	FILE* descriptor returned
a	No	Yes, (append)		FILE* descriptor returned
r+	Yes, from beginning	Yes, from beginning	Error	FILE* descriptor returned.
w+	Yes, from beginning	Yes, From beginning	New file created, FILE* descriptor returned	Delete file contents (overwrite), File* descriptor returned
a+	Yes, from beginning	Yes, from the end (append only)	New file created, FILE* descriptor returned	FILE* descriptor returned
.. b	Binary flag that can be added for binary IO			

What you (probably) want for the assignment

Reading a Text File (Example)

```
FILE *fp;
int c;
int n = 0;
fp=fopen("myfile.txt","r");
if (fp ==NULL) {
    printf("Error opening
file");
} else {
    do {
        c = fgetc (fp);
        if (c == 'A') {
            n++;
        }
    } while (c != EOF);
    fclose (fp);
}
printf("%i", n);
```

Alternatives...

```
int fgetc(FILE *fp);
char *fgets(char *str, int count, FILE *fp);
size_t fread(void *buf, size_t size,
              size_t count, FILE *fp);
```

Writing a Text File (Example)

```
FILE *fp;
char name[];

printf("Enter your name: ");
fgets (name, 256, stdin);

fp=fopen("myfile.txt","a");

if (fp ==NULL) {
    printf("Error opening file");
} else {
    if ( fputs (name, fp) < 0) {
        printf("Error writing file");
    }
    fclose (fp);
}
```

Alternatives...

```
int fputc(int character, FILE *fp);
int *fputs(char *str, FILE *fp);
size_t fwrite(void *ptr, size_t size,
              size_t count, FILE *fp);
```

Always keep in mind


- All those functions will change the state of the file descriptor by advancing the cursor position
- Sometimes not all data requested is read/written
 - Keep track of how many bytes are processed
 - Loop until finished

Reading Formatted strings

- `int fscanf(FILE *stream, const char *format, ...);`
 - Version of `scanf()` that reads from given `FILE*`.
 - `int a; char b[10]; fscanf(stdin, "%d %s", &a, &b);`
 - Also `sscanf(const char *str, const char *fmt, ...)` may be useful
- Format string:
 - `%s` Read string
 - `%d` Read integer

File IO: Resources

- Reference and Examples
<http://www.cplusplus.com/reference/cstdio/>
- Read the man pages!
 - man 3 getc
 - man 3 isspace
 - man 3 scanf
- <http://stackoverflow.com>



Function pointers basics

Write a C program that has a function that:

- accepts a function pointer and an array of integers
- invokes the pointed-to function with each of the elements in the array as an argument
- overrides the current array element with the return value of the called function

```
void map(int (*f) (int), int *array, size_t len) {  
  
}  
  
int pow2(int a) {  
    return a*a;  
}
```

Function pointer

This last part will help you get even more familiar with function pointers.

- write a callback fct to use `qsort()` to sort records based on first name
- write another callback function for last name
- use `apply` and a callback to filter the records
- BONUS: implement your own `mysort()`

Short Quiz about pointers

to get you started thinking of pointers

Quiz: Simple Pointer

```
int a[] = { 0,1,2,3,4 };
```

```
int i, *p;
```

```
for (p = &a[0], i=0; p+i <= a+4; p++,i++)  
    printf("*(p+i) = %d", *(p+i));
```

Solution: Simple Pointer

- $*p+i = 0$ $*p+i = 2$ $*p+i = 4$

Quiz: Arrays and Pointers

```
int a[] = { 0, 1, 2, 3, 4 };  
int *p[] = {a, a+1, a+2, a+3, a+4 };  
int **pp = p;
```

```
main() {  
    printf("...", a, *a, **a);  
    printf("...", p, *p, **p);  
    printf("...", pp, *pp, **pp);  
}
```

Quiz: Arrays and Pointers

- `a` = address of `a`
 `*a` = `0`
 `**a` = Segmentation Fault (Null pointer dereference)
- `p` = address of `p`
 `*p` = address of `a`
 `**p` = `0`
- `pp` = address of `p`
 `*pp` = address of `a`
 `**pp` = `0`

See you
next week!

MAN, I SUCK AT THIS GAME.
CAN YOU GIVE ME
A FEW POINTERS?

0x3A28213A
0x6339392C,
0x7363682E.

I HATE YOU.

