



Chapter 1: Introduction

252-0061-00 V Systems Programming and Computer Architecture

This course covers in *depth*...

- How to write **fast** and **correct** code
- How to write good **systems** code
- What makes programs go fast (and slow)
- Programming in C
 - *Still* the systems programming language of choice
- Programming in Assembly Language
 - What the machine understands
- Programs as more than mathematical objects
 - E.g. how does Facebook work?
 - How programs **interact** with the hardware

Who are we?



Prof. Timothy Roscoe

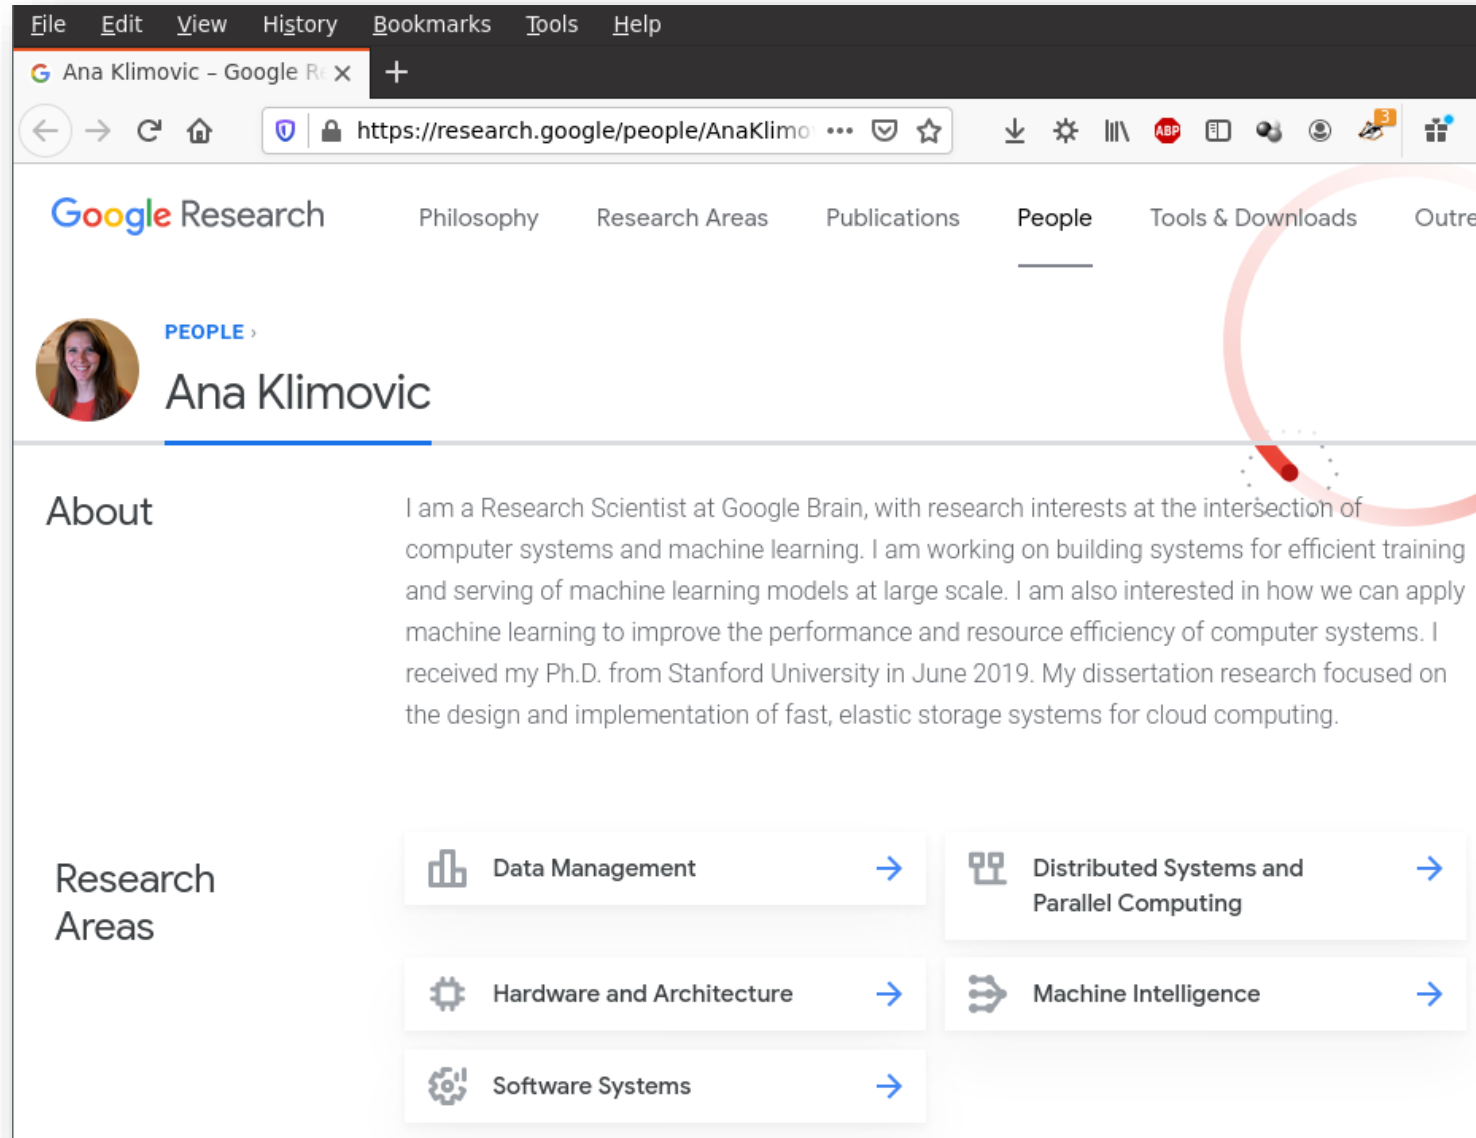


Prof. Ana Klimovic

Full Disclosure



Full Disclosure




The screenshot shows a web browser window with the Google Research profile of Ana Klimovic. The browser's address bar shows the URL `https://research.google/people/AnaKlimovic`. The page has a navigation bar with links to Philosophy, Research Areas, Publications, People (which is underlined), Tools & Downloads, and Outreach. Below the navigation bar, there is a profile picture of Ana Klimovic, a blue 'PEOPLE' label, and her name 'Ana Klimovic'. The 'About' section contains a paragraph about her role as a Research Scientist at Google Brain and her research interests. Below this, the 'Research Areas' section lists five categories: Data Management, Distributed Systems and Parallel Computing, Hardware and Architecture, Machine Intelligence, and Software Systems, each with a corresponding icon and a right-pointing arrow. A red circle is drawn around the 'About' section.

File Edit View History Bookmarks Tools Help

Ana Klimovic - Google Research






Google Research Philosophy Research Areas Publications People Tools & Downloads Outreach

 PEOPLE Ana Klimovic

About

I am a Research Scientist at Google Brain, with research interests at the intersection of computer systems and machine learning. I am working on building systems for efficient training and serving of machine learning models at large scale. I am also interested in how we can apply machine learning to improve the performance and resource efficiency of computer systems. I received my Ph.D. from Stanford University in June 2019. My dissertation research focused on the design and implementation of fast, elastic storage systems for cloud computing.

Research Areas

-  Data Management →
-  Distributed Systems and Parallel Computing →
-  Hardware and Architecture →
-  Machine Intelligence →
-  Software Systems →

1.1: Logistics

Systems Programming and Computer Architecture

Lectures

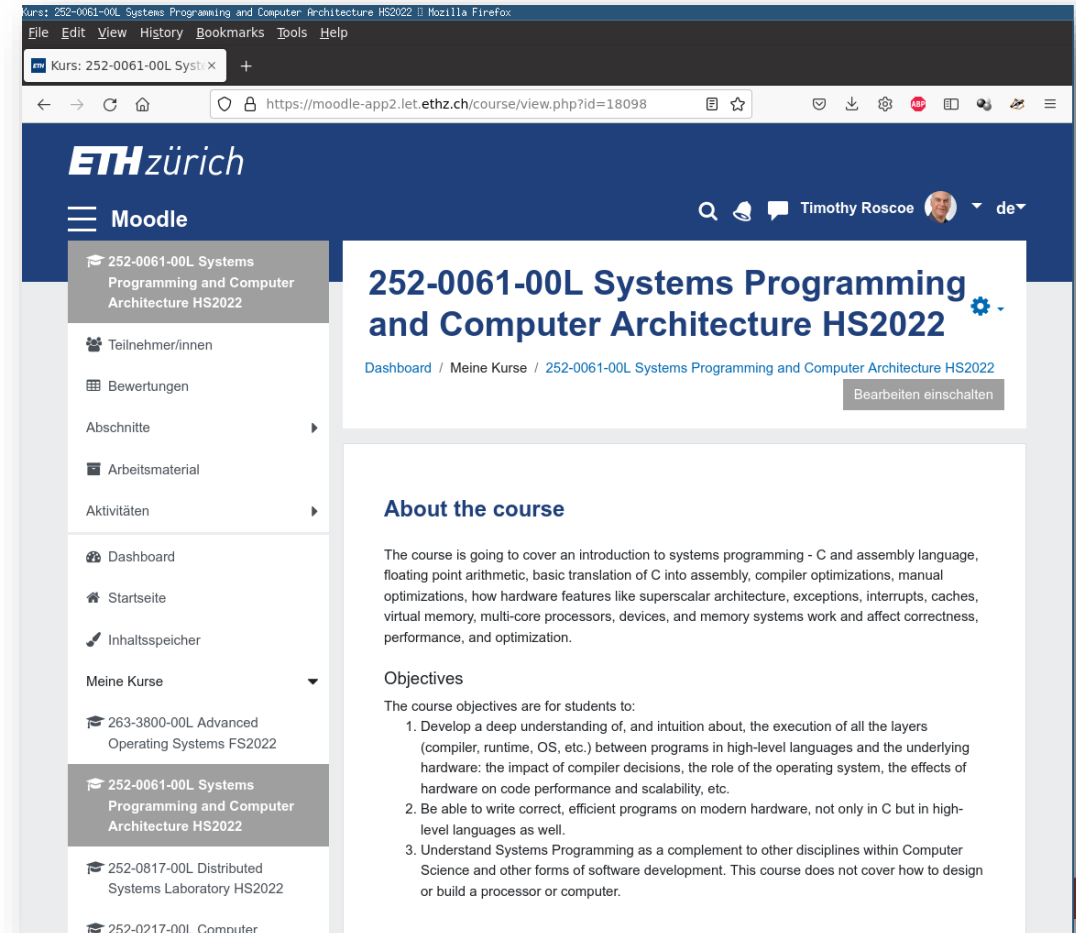
*The slides are **not**
intended to be
understood without
the lectures...*

- Physical:
 - 10:00-12:00 Tuesdays and Wednesdays
 - HG E 7 (Tue) & NO C 60 (Wed)
- Recordings will appear after a few days
 - <https://video.ethz.ch/lectures/d-infk/2021/autumn.html>

Moodle

<https://moodle-app2.let.ethz.ch/course/view.php?id=18098>

- The first place to look!
- **Links** posted here
- All **lecture materials** will be posted on Moodle.
- Ask **questions** in the **forum**
 - TAs and Profs will monitor the forum!
- We will **not** answer questions on Discord.



Tutorial sessions

- Very important!
- Logistics:
 - Wednesday, 12:00-14:00 or 14:00-16:00
 - See myStudies for rooms and **streams**
- Content:
 - Tools and skills for **lab exercises**
 - Knowledge needed for exams, but not in the lectures!
- There **will** be a session this Wednesday (tomorrow)

Language

- We'll teach in English (and C...)
 - If we speak too fast, or say something unclear, raise your hand!
 - Please ask questions!
- Assistants speak German, English, Italian, French, ...
- Examination (100% of grade):
 - Paper will be in English
 - Answers may be in German or English


Asking questions

1. Ask **during the lectures**
2. Ask on the **Moodle forum** outside the lectures
3. Ask your **friends**
4. Check the **web**
5. Ask your teaching **assistant**
6. Ask ***another*** teaching assistant
7. Email us (troscoe@inf.ethz.ch or aklimovic@ethz.ch)

Acknowledgements

- Lots of material from the famous **CS 15-213** at Carnegie Mellon University
 - Basis for the **book**
- Some C programming slides adapted from **CSE333** at University of Washington
 - Many thanks to (ex-)Prof. Steve Gribble
- New material:
 - Considerable evolution...
 - Multicore, devices, etc.
 - Mostly Roscoe's fault 😊

Questions?



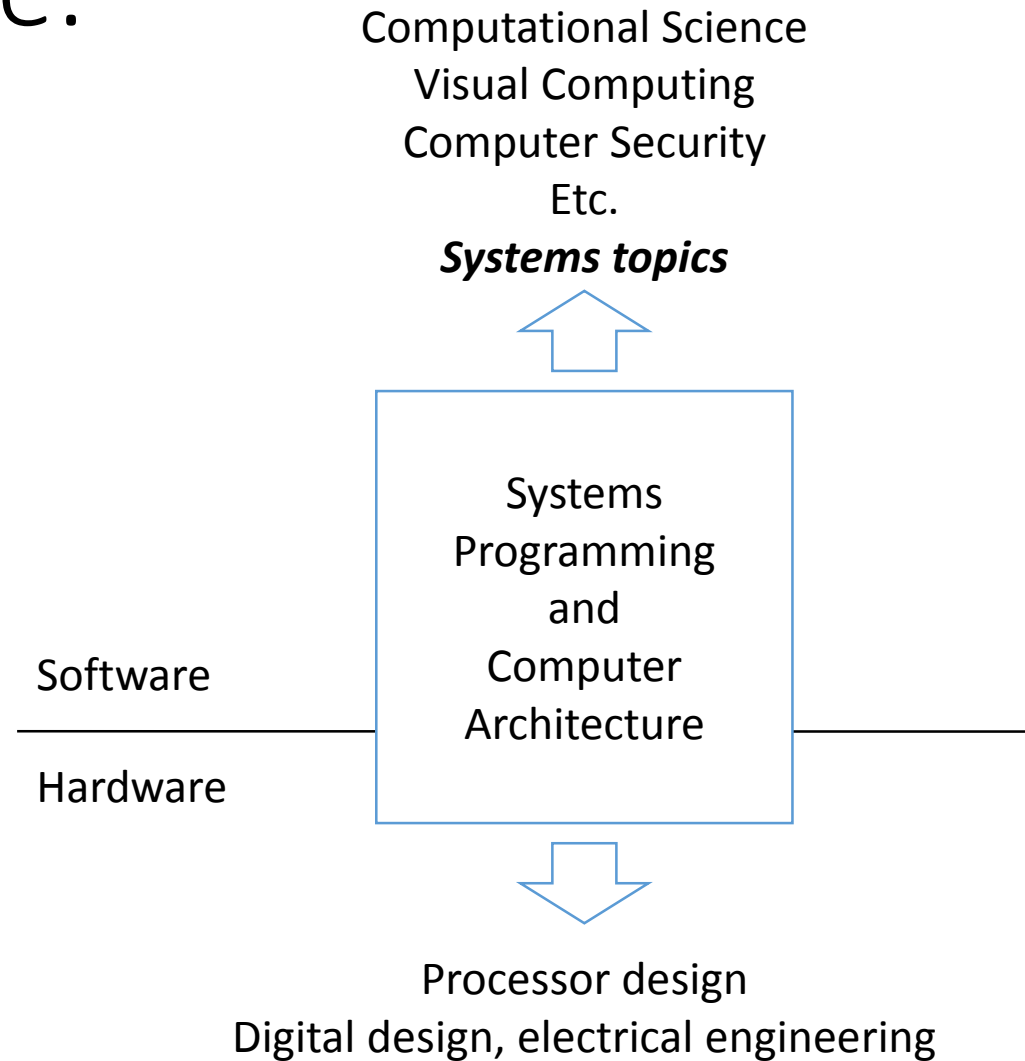
1.2: What is Systems Programming?

Systems Programming and Computer Architecture

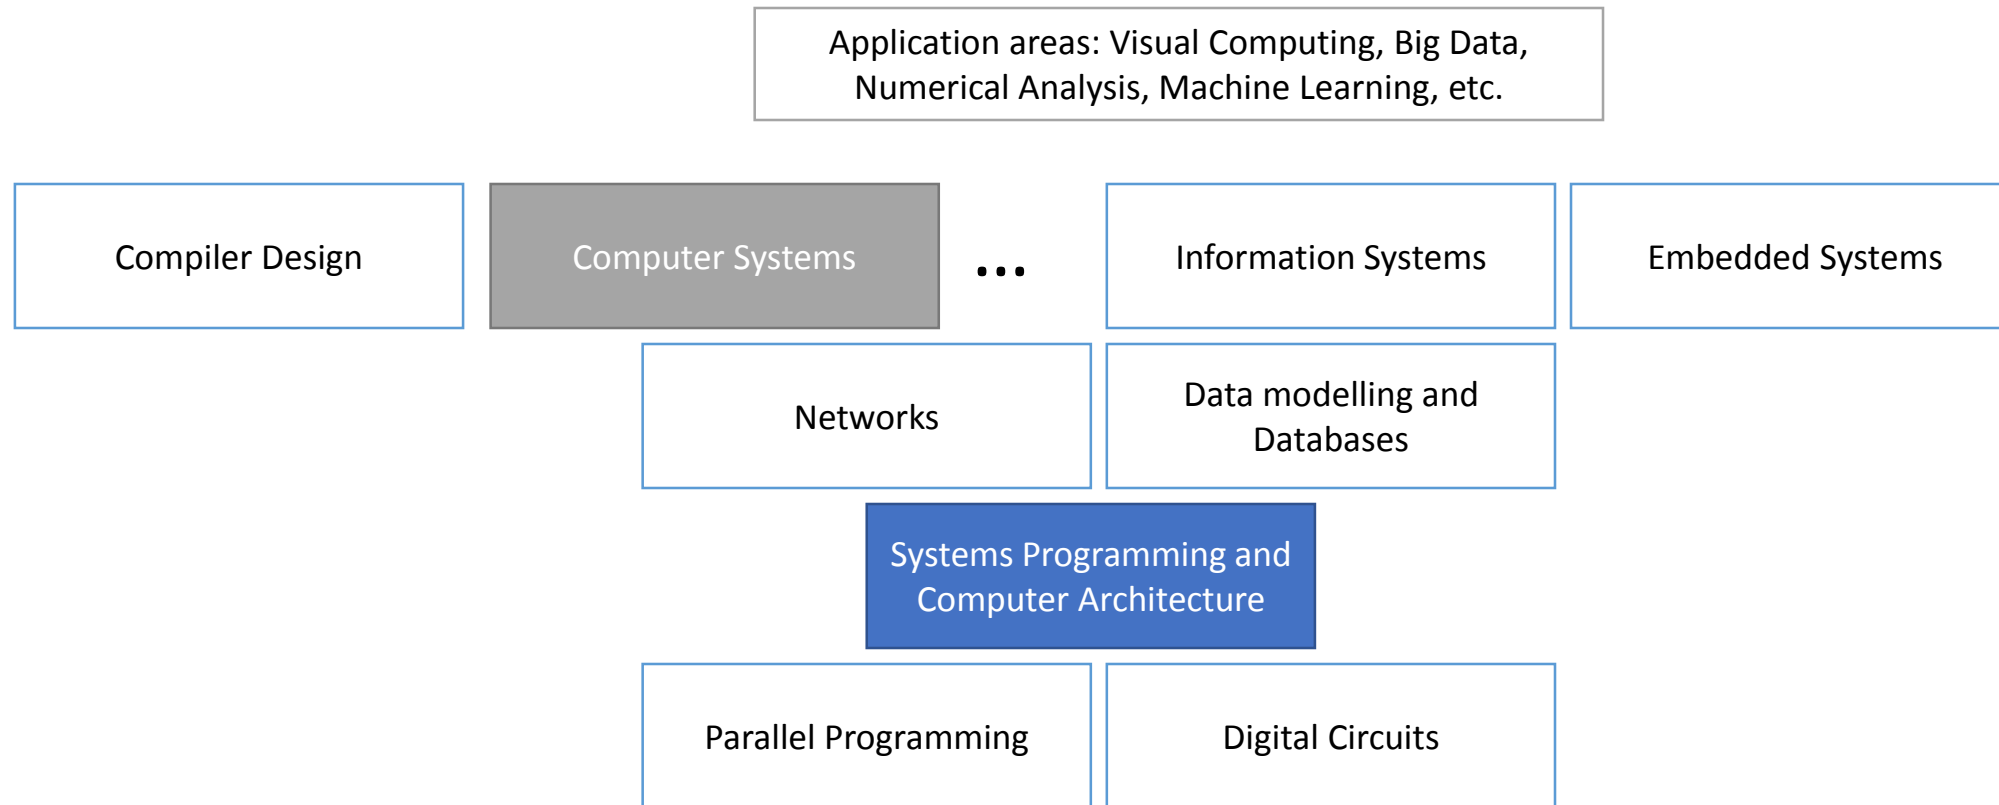
“Systems” as a field

- Encompasses:
 - Operating systems
 - Database systems
 - Networking protocols and routing
 - Compiler design and implementation
 - Distributed systems
 - Cloud computing & online services
 - Big Data and machine learning frameworks
- On and above the *hardware/software* boundary

You are here:



You are here:



Motivation

- Most CS courses emphasize **abstraction**
 - Abstract data types (objects, contracts, etc.)
 - Asymptotic analysis (worst-case, complexity)
- These abstractions have **limitations**
 - Often don't survive contact with reality
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations

Summary: Course Goals

- Become more effective programmers
 - Find and eliminate **bugs** efficiently
 - Understand and tune for program **performance**
- Prepare for later **systems** classes at ETHZ
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems

Questions?

1.3: Motivation - Six realities


Systems Programming and Computer Architecture

Reality #1:

Computers don't really deal
with numbers.

ints are not integers, floats are not reals

- Is $x^2 \geq 0$?
 - floats: Yes!
 - ints:
 - $40000 * 40000 \rightarrow 1600000000$
 - $50000 * 50000 \rightarrow ??$
- Is $(x + y) + z = x + (y + z)$?
 - unsigned & signed ints: Yes!
 - floats:
 - $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
 - $1e20 + (-1e20 + 3.14) \rightarrow ??$



<http://xkcd.com/571>

Computer arithmetic

- Does not generate random values
 - Arithmetic operations have important mathematical properties
- Cannot assume all “usual” mathematical properties
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs
- Observation
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programmers

Reality #2:

You've got to know assembly.

You've got to know assembly

- Chances are, you'll never write a program in assembly
 - Compilers are much better & more patient than you are
- But: understanding assembly is **key** to machine-level execution model
 - Behavior of programs in presence of **bugs**
 - High-level language model breaks down
 - Tuning program **performance**
 - Understand optimizations done/not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing **system software**
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting **malware**
 - x86 assembly is the language of choice!

Assembly code example

- Time Stamp Counter
 - Special 64-bit register in Intel-compatible machines
 - Incremented every clock cycle
 - Read with `rdtsc` instruction
- Application
 - Measure time (in clock cycles) required by procedure

```
double t;  
start_counter();  
P();  
t = get_counter();  
printf("P required %f clock cycles\n", t);
```

Code to read counter

- Write small amount of assembly code using GCC's asm facility
- Inserts assembly code into machine code generated by compiler

```
uint64_t access_counter()
{
    uint32_t lo, hi;

    asm volatile("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (hi), "=r" (lo)
        :
        : "%edx", "%eax");

    return (lo | (((uint64_t)hi) << 32));
}
```

Reality #3:

Memory matters.

RAM is not a realistic abstraction.

Memory matters

- Memory is **not unbounded**
 - It must be allocated and managed
 - Many applications are memory-dominated
- Memory performance is **not uniform**
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements
- Memory is **typed**
 - Different kinds of memory behave differently

... and memory-related bugs are a nightmare.

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

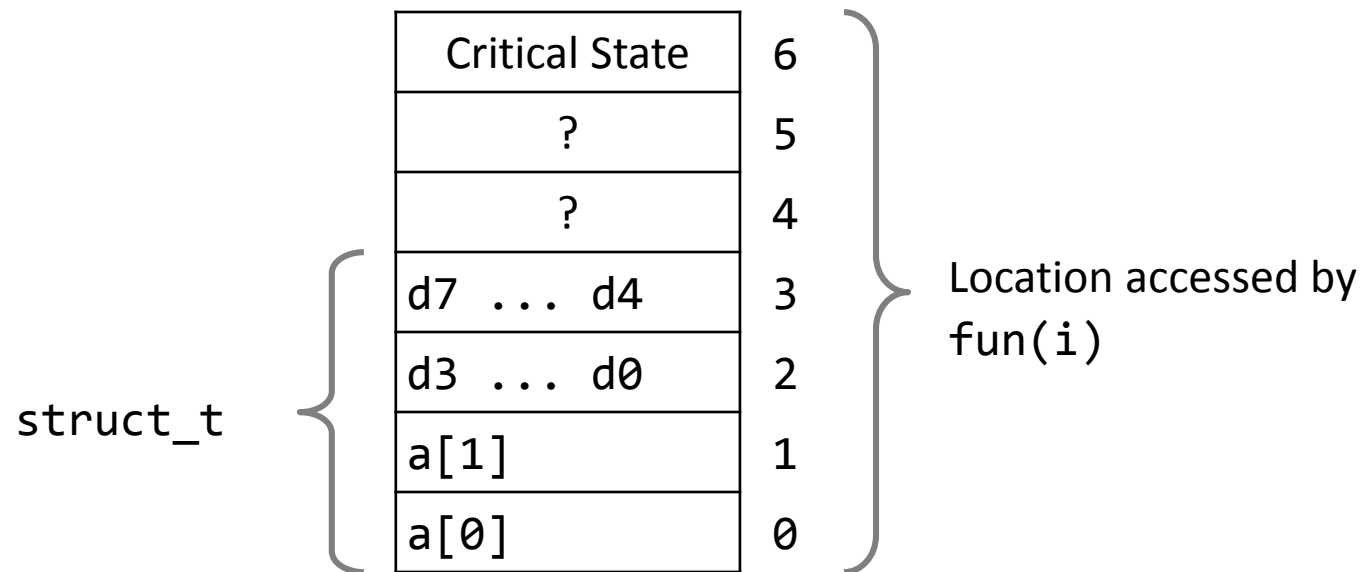
Actual results are
system-specific...

Memory referencing bug

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

Explanation:



Memory system performance

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

5.2 ms


```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

162 ms !

Intel Core i7 2.7 GHz

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

The Memory Mountain



Reality #4:

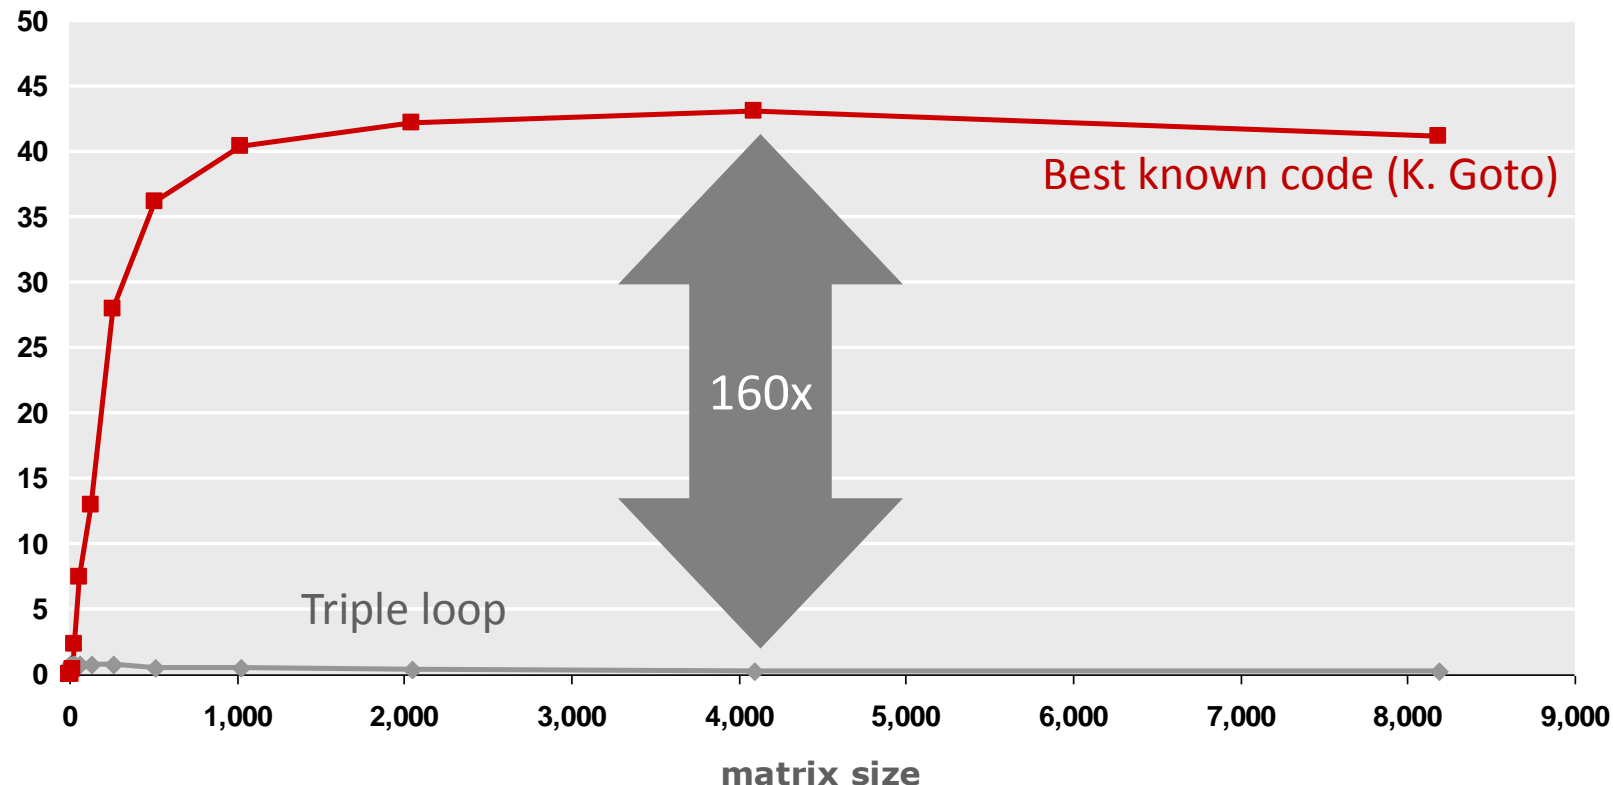
There's much more to performance than
asymptotic complexity

There's much more to performance than asymptotic complexity

- Constant factors matter too!
- Even exact op count does not predict performance
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example: matrix-matrix multiplication


Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)
Gflop/s



- Standard desktop computer, vendor compiler, using optimization flags
- Both implementations have **exactly** the same operations count ($2n^3$)
- What is going on?

MMM plot: analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz
Gflop/s



- Why? Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, ...
- **Effect: less register spills, less L1/L2 cache misses, less TLB misses**

Reality #5:

Computers don't just execute programs
Programs don't just calculate values

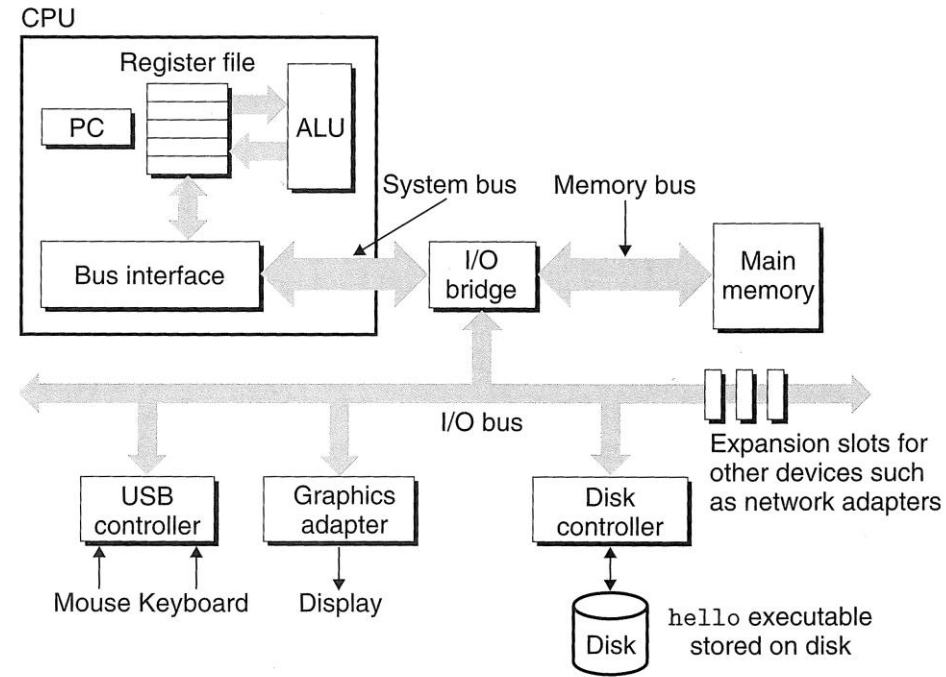
Computers don't just run programs

- They need to get data **in** and **out**
 - I/O critical to program reliability and performance
 - **Sense** the physical world
 - **Act** in the physical world
- They **communicate** over networks
 - Many system-level issues arise with a network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross-platform interoperability
 - Complex performance issues

Lies our teachers tell us...

Figure 1.4

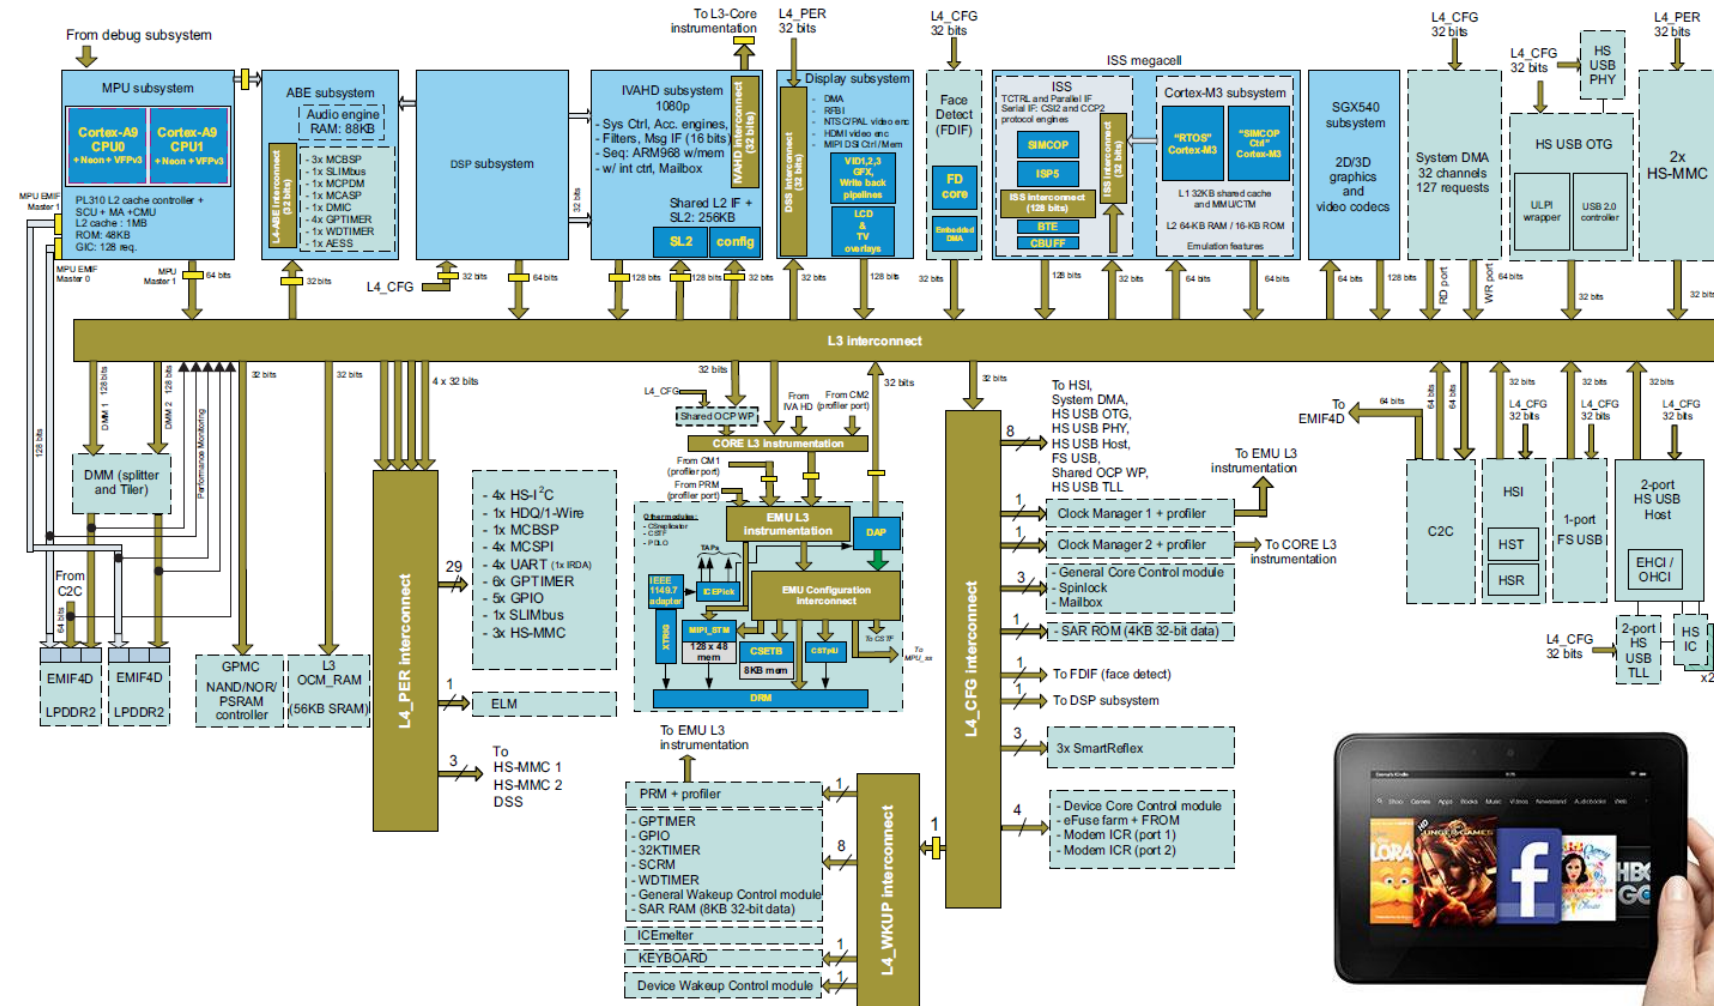
Hardware organization of a typical system. CPU: Central Processing Unit, ALU: Arithmetic/Logic Unit, PC: Program counter, USB: Universal Serial Bus.



systems, but all systems have a similar look and feel. Don't worry about the complexity of this figure just now. We will get to its various details in stages throughout the course of the book.

Computer Systems, A Programmer's Perspective, Bryant & O'Hallaron, 2011

A modern(ish) System-on-Chip



Texas Instruments
OMAP 4460, c.2011

Reality #6:

Programs are not complete semantic specifications

The role of “standards”

- Language standards aim to **specify unambiguously** what any program in the language does when compiled and executed.
 - Java: “write once, run anywhere”
 - Formal semantics
- The C standards should be viewed as rather different
 - Behavior frequently described as “**implementation dependent**”
 - What does this mean?

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two options:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

access

⟨execution-time action⟩ to read or modify the value of an object

NOTE 1 Where only one of these two actions is meant, “read” or “modify” is used.

NOTE 2 “Modify” includes the case where the new value being stored is the same as the previous value.

NOTE 3 Expressions that are not evaluated do not access objects.

3.2

alignment

requirement that objects of a particular type be located on storage boundaries with addresses that are particular multiples of a byte address

3.3

argument

actual argument

actual parameter (deprecated)

expression in the comma-separated list bounded by the parentheses in a function call expression, or a sequence of preprocessing tokens in the comma-separated list bounded by the parentheses in a function-like macro invocation

3.4

behavior

external appearance or action

3.4.1

implementation-defined behavior

unspecified behavior where each implementation documents how the choice is made

EXAMPLE An example of implementation-defined behavior is the propagation of the high-order bit when a signed integer is shifted right.

3.4.2

locale-specific behavior

behavior that depends on local conventions of nationality, culture, and language that each implementation documents

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two interpretations:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

Default behavior for newer C compilers ☹️

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

...ect
...dify” is used.
... is the same as the previous value.
... on storage boundaries with

argument

actual argument
actual parameter (deprecated)
expression in the comma-separated list bounded by the parentheses in a function call expression, or a sequence of preprocessing tokens in the comma-separated list bounded by the parentheses in a function-like macro invocation

3.4

behavior

external appearance or action

3.4.1

implementation-defined behavior

unspecified behavior where each implementation documents how the choice is made

EXAMPLE An example of implementation-defined behavior is the propagation of the high-order bit when a signed integer is shifted right.

3.4.2

locale-specific behavior

behavior that depends on local conventions of nationality, culture, and language that each implementation documents

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two interpretations:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

Default behavior for newer C compilers ☹️

Default behavior for older C compilers, and *What You Actually Want.*

The role of “standards”

- Language standards aim to **specify unambiguously** what any program in the language does when compiled and executed.
 - Java: “write once, run anywhere”
 - Formal semantics
- The C standards should be viewed as rather different
 - Behavior frequently described as “**implementation dependent**”
 - What does this mean?


A program is a set of instructions to a compiler that tell it what assembly language to generate.

Summary

1. ints are not integers. floats are not real numbers.
2. You've got to know assembly.
3. Memory matters. RAM is an unrealistic abstraction.
4. There's much more to performance than asymptotic complexity.
5. Computers don't just evaluate programs.
6. Programs are not complete semantic specifications.


Interesting times...

- Processors are not getting faster.
- Performance-wise, progress in computer architecture has halted.
 - In fact, it's going backwards due to security concerns.




Interesting times...

- Lots of companies, universities, etc. are building new kinds of computers.
- How can these be programmed?
- This is a systems software problem.



Computers are looking different!

- Cavium / Marvell ThunderX-1 NP:
 - 48 x ARMv8 cores at 2GHz
 - 128 GB DDR4
 - 2 x 40Gb/s network
- Xilinx UltraScale+ VU9P
 - 512 GB DDR4
 - 4 x 100 Gb/s network
- NVMe, SATA, PCIe on both sides
- Native coherence between FPGA and CPU



1.4: What we'll assume you know

Systems Programming and Computer Architecture

Courses already

- Programming & software engineering
- Parallel programming
- Data structures and algorithms
- Digital Circuits
- Discrete Mathematics

What we'll assume you know #1:


- Binary, and Hexadecimal notation
- Memory, addresses, bytes, and **words**
- Byte-ordering
(Big-Endian, Little-Endian)
- Boolean algebra
(and, or, not, xor)
- Generalized Boolean algebra
(bitwise operations on words as bit vectors)
- How to write programs
- Some languages
 - Java (but we won't use it)
 - A bit of C (but not much)
- Some MIPS (or other) assembly

What we'll assume you know #2:

- Processor architecture, pipelines
- MIPS, ARM, or LC-3b assembly (we'll work in 64-bit x86)
 - Registers
 - Addressing modes
 - Instruction formats
- Basic memory systems
 - cache architectures
 - virtual memory
 - I/O devices
- Software engineering
 - Object-orientation
 - Design-by-contract
 - Strong typing
- Concurrency and parallelism
 - Threads
 - Locks, mutexes, condition variables
 - Parallel programming constructs



Textbooks

- Randal E. Bryant and David R. O'Hallaron,
 - Computer Systems: A Programmer's Perspective, Third Edition (CS:APP3e), Pearson, 2016
 - <http://csapp.cs.cmu.edu>




Books on C (there are many)

- Brian Kernighan and Dennis Ritchie,
 - “The C Programming Language, Second Edition”, Prentice Hall, 1988(!)



- Samuel Harbison and Guy Steele
 - C: A Reference Manual
 - 5th edition 2002

- Peter van der Linden
 - Expert C Programming: Deep C Secrets, 1994



OK!