



Chapter 3: Representing integers in C

252-0061-00 V Systems Programming and Computer Architecture

Goal:

Introduction to integer arithmetic on computers

- Recap of how numbers are represented
- Signed and unsigned values
- Integer ranges
- Integer addition and subtraction
 - In C
 - Mathematical properties
- Integer multiplication
 - In C
 - Mathematical properties
- Integer multiplication and division using shifts

3.1: Recap: Encodings and operators

Computer Architecture and Systems Programming

Representing integers


```
int A = 15213;  
int B = -15213;  
long int C = 15213;
```

Decimal: 15213


Binary: 0011 1011 0110 1101

Hex: 3 B 6 D


A on ia32, x86-64




A on SPARC




C on ia32




C on x86-64




C on SPARC



B on ia32, x86-64



B on SPARC



Two's complement representation

Bit-level operations in C

- Operations `&`, `|`, `~`, `^` available in C
 - Apply to any “integral” data type
 - `long`, `int`, `short`, `char`, `unsigned`
 - View arguments as bit vectors
 - Arguments applied bit-wise

- Examples (using `char` data type):

- $\sim 0x41 \rightarrow 0xBE$
 $\sim 01000001_2 \rightarrow 10111110_2$
- $\sim 0x00 \rightarrow 0xFF$
 $\sim 00000000_2 \rightarrow 11111111_2$
- $0x69 \ \& \ 0x55 \rightarrow 0x41$
 $01101001_2 \ \& \ 01010101_2 \rightarrow 01000001_2$
- $0x69 \ | \ 0x55 \rightarrow 0x7D$
 $01101001_2 \ | \ 01010101_2 \rightarrow 01111101_2$

<code>&</code>	Bitwise AND
<code> </code>	Bitwise OR
<code>~</code>	Bitwise NOT
<code>^</code>	Bitwise XOR

Contrast: Logic operations in C

- `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - **Early termination**
- Examples (char data type)
 - `!0x41` \rightarrow `0x00`
 - `!0x00` \rightarrow `0x01`
 - `!!0x41` \rightarrow `0x01`

 - `0x69 && 0x55` \rightarrow `0x01`
 - `0x69 || 0x55` \rightarrow `0x01`

<code>&&</code>	Logical AND
<code> </code>	Logical OR
<code>!</code>	Logical NOT

Representing & manipulating sets

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$
- $a_j = 1$ if $j \in A$:

01101001 $\{0, 3, 5, 6\}$
76543210

01010101 $\{0, 2, 4, 6\}$
76543210

Operator	Operation	Result	Meaning
&	Intersection	01000001	$\{0, 6\}$
	Union	01111101	$\{2, 3, 4, 5, 6\}$
^	Symmetric difference	00111100	$\{2, 3, 4, 5\}$
~	Complement	10101010	$\{1, 3, 5, 7\}$

Shift operations

- Left shift: $x \ll y$
 - Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right shift: $x \gg y$
 - Shift bit-vector x right y positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on right
- **Undefined** behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010 000
Log. $\gg 2$	00 011000
Arith. $\gg 2$	00 011000

Argument x	10100010
$\ll 3$	00010 000
Log. $\gg 2$	00 101000
Arith. $\gg 2$	11 101000

Java writes
this “ \ggg ”.

Summary

- Integer encoding
- Bit-level operators
- Set representation
- Logic operators
- Shift operators

3.2: Integer ranges

Computer Architecture and Systems Programming

Encoding integers

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

```
short int x = 15213;  
short int y = -15213;
```

Sign
Bit

- A C short is 2 bytes long:

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011

- Sign bit
 - For 2's complement, most significant bit = 1 indicates negative

Numeric ranges

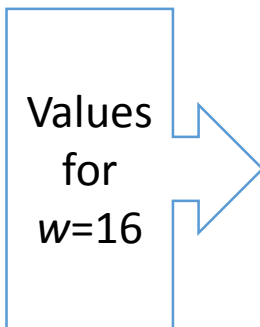
- Unsigned values

- UMin = 0
 - 000...0
- UMax = $2^w - 1$
 - 111...1

- Two's complement values

- TMin = -2^{w-1}
 - 100...0
- TMax = $2^{w-1} - 1$
 - 011...1

Values
for
 $w=16$



	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for different word sizes

	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

- Observations:

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$

- C Programming

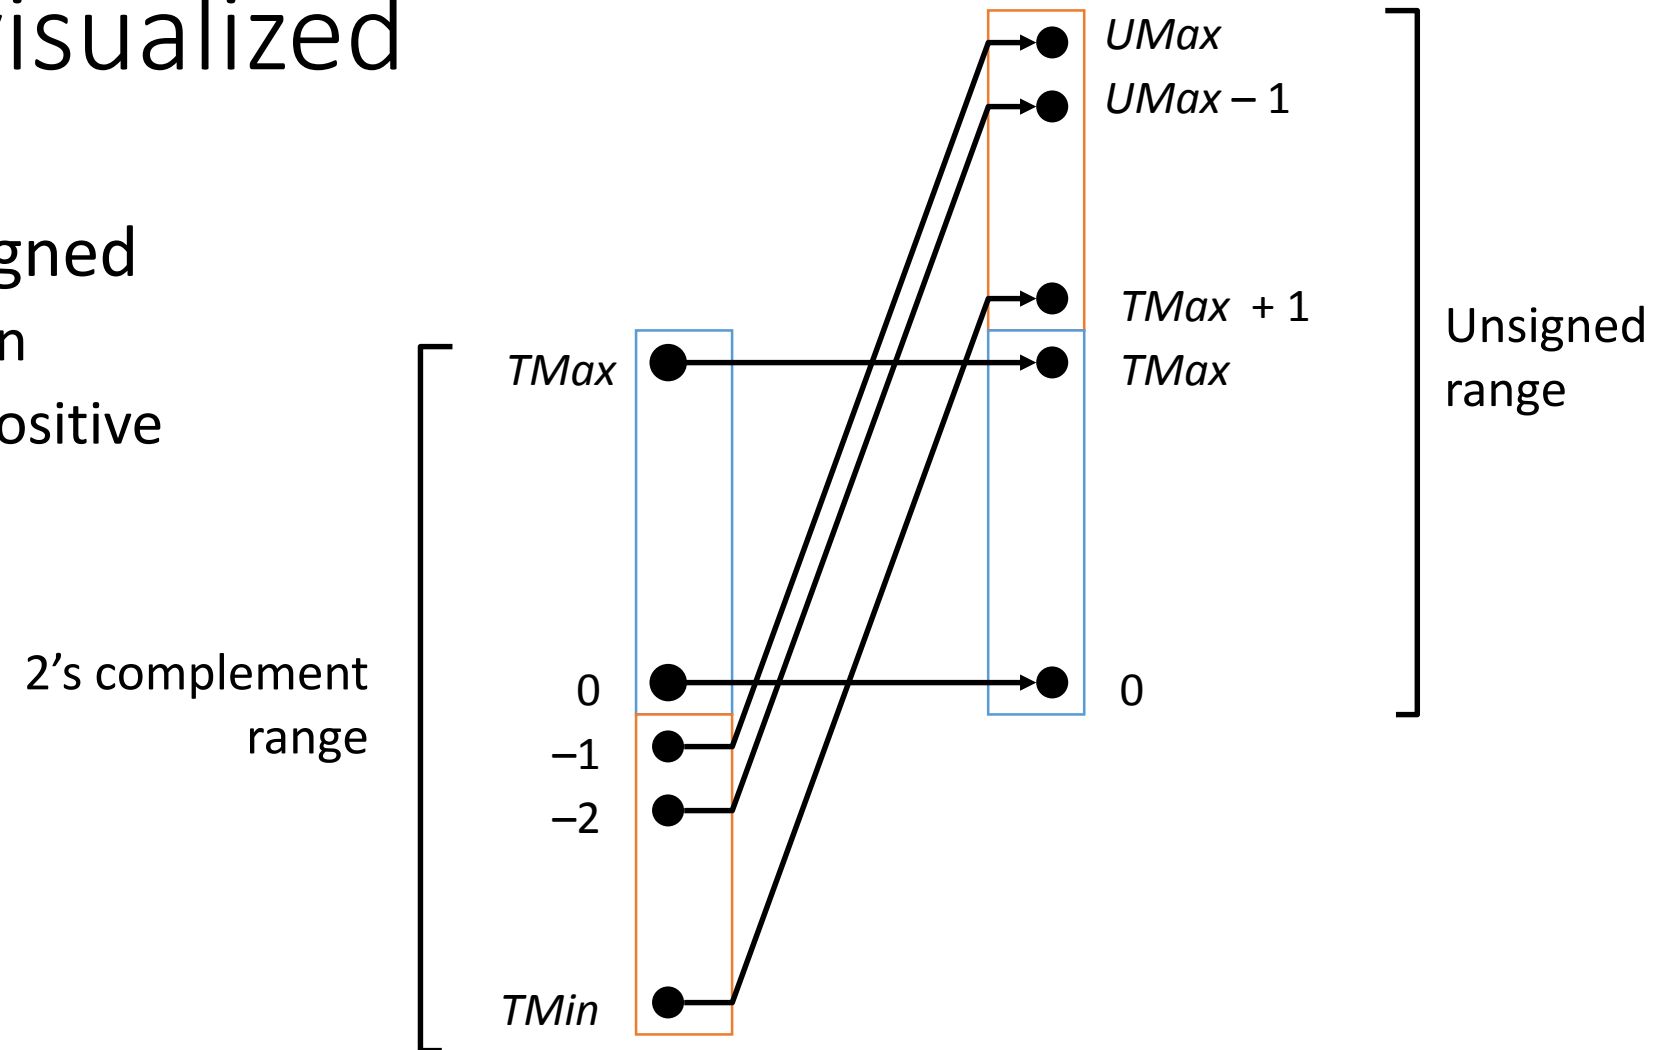
- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values platform specific

Mapping signed \leftrightarrow unsigned

Bits	Signed		Unsigned
0000	0	\longleftrightarrow =	0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8	\longleftrightarrow +16	8
1001	-7		9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15

Conversion visualized

- 2's Comp. \rightarrow Unsigned
 - Ordering inversion
 - Negative \rightarrow big positive



Signed vs. unsigned in C

- Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix:

```
0U  
4294967259U
```

- Casting

- Explicit casting between signed & unsigned same as U2T and T2U:

```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

- Implicit casting also occurs via assignments and procedure calls

```
tx = ux;  
uy = ty;
```


Casting surprises in expression evaluation

- When mixing unsigned and signed in single expression, *signed values implicitly cast to unsigned*
 - Including comparison operators!
<, >, ==, <=, >=
- Examples for W = 32:
 - TMIN = -2,147,483,648 ,
 - TMAX = 2,147,483,647

Constant 1	Constant 2	Relation	Evaluation
0	0U	==	Unsigned
-1	0	<	Signed
-1	0U	>	Unsigned
2147483647	-2147483647-1	>	Signed
2147483647U	-2147483647-1	<	Unsigned
-1	-2	>	Signed
(unsigned)-1	-2	>	Unsigned
2147483647	2147483648U	<	Unsigned
2147483647	(int) 2147483648U	>	signed

Code security example from way back (but...)

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- Similar to code found in FreeBSD's implementation of `getpeername`
- There are legions of smart people trying to find vulnerabilities in programs...

Malicious usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* Declaration of library function memcpy */
void *memcpy(void *dest, void *src, size_t n);


#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```

Sign extension

- Task:
 - Given w -bit signed integer x
 - Convert it to $w+k$ -bit integer with same value

- Rule:
 - Make k copies of sign bit:
 - $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$



Sign extension example

- Converting from smaller to larger integer data type
- C automatically performs sign extension for signed values

```
short int x = 15213;  
int      ix = (int) x;  
short int y = -15213;  
int      iy = (int) y;
```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
ix	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
y	-15213	C4 93	11000100 10010011
iy	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

Summary

- **Casting** signed \leftrightarrow unsigned
 - Bit pattern is maintained but **reinterpreted**
 - Can have unexpected effects: adding or subtracting 2^{w-1}
- Expression containing signed and unsigned int
 - int is **cast to unsigned**!
- Expanding (e.g., short int \rightarrow int)
 - Unsigned: zeros added
 - Signed: **sign extension**
- Truncating (e.g., unsigned \rightarrow unsigned short)
 - Unsigned: modulus operation
 - Signed: similar to modulus
 - For small numbers yields expected behaviour

3.3: Integer addition and subtraction in C

Computer Architecture and Systems Programming

Negation: complement & increment

- Recall the following holds for 2's complement:

$$\sim x + 1 == -x$$

- Complement

Observation: $\sim x + x == 1111\dots111 == -1$

x	1	0	0	1	1	1	0	1
+ ~x	0	1	1	0	0	0	1	0
<hr/>								
-1	1	1	1	1	1	1	1	1

- Complete proof?

Complement & increment examples

$x = 15213$

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
$\sim x$	-15214	C4 92	11000100 10010010
$\sim x + 1$	-15213	C4 93	11000100 10010011
y	-15213	C4 93	11000100 10010011

$x = 0$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~ 0	-1	FF FF	11111111 11111111
$\sim 0 + 1$	0	00 00	00000000 00000000

Unsigned addition

Operands: w bits



True sum: $w+1$ bits



Discard carry: w bits




- Standard addition function
 - Ignores carry output
- Implements modular arithmetic

$$s = UAdd_w(u, v) = u + v \bmod 2^w$$

$$UAdd_w(u, v) = \begin{cases} u + v, & u + v < 2^w \\ u + v - 2^w, & u + v \geq 2^w \end{cases}$$



Visualizing (mathematical) integer addition

- 4-bit integers u, v
- Compute true sum $\text{Add}_4(u, v)$
- Values increase linearly with u and v
- Forms planar surface



Visualizing **unsigned** addition

- Wraps around
 - If true sum $\geq 2^w$
 - At most once



Mathematical properties

- Modular addition forms an Abelian group

- **Closed** under addition

$$0 \leq UAdd_w(u, v) \leq 2^w - 1$$

- **Commutative**

$$UAdd_w(u, v) = UAdd_w(v, u)$$

- **Associative**

$$UAdd_w(t, UAdd_w(u, v)) = UAdd_w(UAdd_w(t, u), v)$$

- 0 is additive **identity**

$$UAdd_w(u, 0) = u$$

- Every element has additive **inverse**

Let: $UComp_w(u) = 2^w - u$

Then: $UAdd_w(u, UComp_w(u)) = 0$

Two's complement addition

Operands: w bits

u 

$+ v$ 

True sum: $w+1$ bits

$u + v$ 

Discard carry: w bits

$TAdd_w(u, v)$ 

- TAdd and UAdd have identical bit-level behavior
 - Signed vs. unsigned addition in C:


```
int s, t, u, v;  
s = (int) ((unsigned) u + (unsigned) v);  
t = u + v;
```

will give:

$s == t$

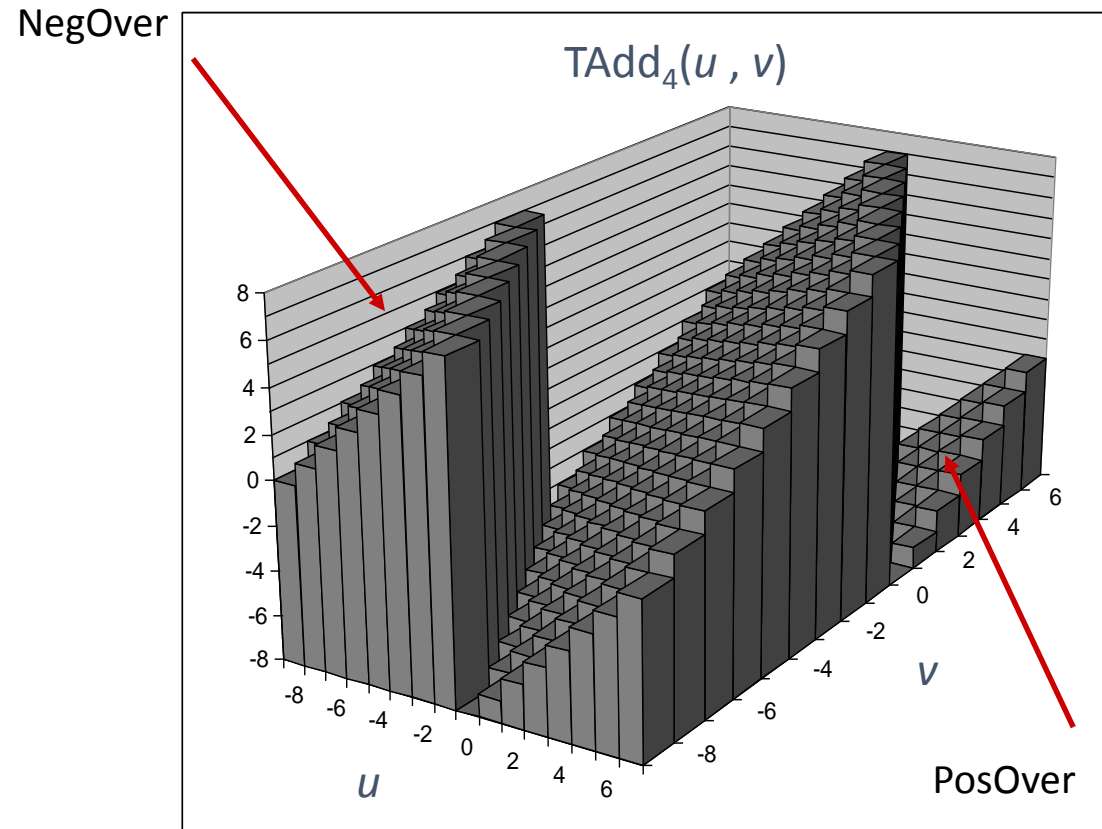
TAdd overflow

- Functionality
 - True sum requires $w+1$ bits
 - Drop off MSB
 - Treat remaining bits as 2's comp. integer



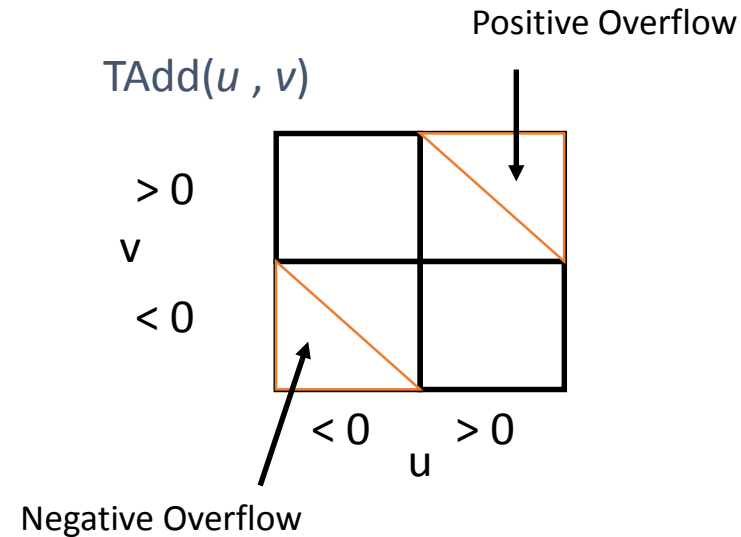
Visualizing 2's complement addition

- Values
 - 4-bit two's comp.
 - Range from -8 to +7
- Wraps around
 - If $\text{sum} \geq 2^{w-1}$
 - Becomes negative
 - At most once
 - If $\text{sum} < -2^{w-1}$
 - Becomes positive
 - At most once



Characterizing TAdd

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



$$TAdd_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \end{cases}$$

(Neg. overflow)

(Pos. overflow)

Mathematical properties of TAdd

- Group isomorphic to unsigneds with UAdd
 - Since both have identical bit patterns

$$TAdd_w(u, v) = U2T \left(UAdd_w(T2U(u), T2U(v)) \right)$$

- 2's complement under TAdd forms a group
 - Closed, commutative, associative,
 - 0 is additive identity
 - Every element has additive inverse

$$TComp_w(u) = \begin{cases} -u & u \neq TMin_w \\ TMin_w & u = TMin_w \end{cases}$$


3.4: Integer multiplication in C

Computer Architecture and Systems Programming

Multiplication

- Computing exact product of w -bit numbers x, y
 - Either signed or unsigned
- Ranges
 - Unsigned (up to $2W$ bits):
$$0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$$
 - Two's complement min (up to $2W - 1$ bits):
$$x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$$
 - Two's complement max (up to $2W - 2$ bits, but only for $(TMin_w)^2$):
$$x * y \leq (-2^{w-1})^2 = 2^{2w-2}$$
- Maintaining exact results
 - Would need to keep expanding word size with each product computed
 - Done in software by “arbitrary precision” arithmetic packages

Unsigned multiplication in C



- Standard multiplication function
 - Ignores high order w bits
- Implements modular arithmetic

$$UMult_w(u, v) = u \cdot v \bmod 2^w$$

Unsigned multiplication with addition forms a **commutative ring**:

- Addition is a commutative group

- Closed under **multiplication**

$$0 \leq \text{UMult}_w(u, v) \leq 2^w - 1$$

- Multiplication is **commutative**

$$\text{UMult}_w(u, v) = \text{UMult}_w(v, u)$$

- Multiplication is **associative**

$$\text{UMult}_w(t, \text{UMult}_w(u, v)) = \text{UMult}_w(\text{UMult}_w(t, u), v)$$

- 1 is multiplicative **identity**


$$\text{UMult}_w(u, 1) = u$$

- Multiplication **distributes** over addition

$$\text{UMult}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UMult}_w(t, u), \text{UMult}_w(t, v))$$

Signed multiplication in C

Operands: w bits



True product: $2*w$ bits



Discard w bits: w bits



- Standard multiplication function
 - Ignores high order w bits
 - Some of which are different for signed vs. unsigned multiplication
 - Lower bits are the same

Signed multiplication

- Isomorphic algebra to unsigned multiplication and addition
 - Both isomorphic to ring of integers mod 2^w

- Comparison to (mathematical) integer arithmetic

- Both are rings
- True integers obey ordering properties, e.g.,

$$u > 0 \Rightarrow u + v > v$$
$$u > 0, v > 0 \Rightarrow u \cdot v > 0$$

- *Not the case* for two's complement arithmetic:

$$\text{TMax} + 1 == \text{TMin}$$

$$15213 * 30426 == -10030 \quad (\text{e.g. for 16-bit words})$$

3.5: Integer multiplication and division using shifts


Computer Architecture and Systems Programming

Power-of-2 multiply with shift

- Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned


Operands: w bits



True product: $w+k$ bits



Discard k bits: w bits



- Examples

- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Compiled multiplication code

C Function

```
int mul12(int x)
{
    return x*12;
}
```

Compiled arithmetic operations

```
leal  (%eax,%eax,2), %eax
sall  $2, %eax
```

Explanation

```
t <- x+x*2
return t << 2;
```

- C compiler automatically generates shift/add code when multiplying by constant

Unsigned power-of-2 divide with shift

- Quotient of unsigned by power of 2


- $u \gg k$ gives $\lfloor u / 2^k \rfloor$

- Uses logical shift

Operands:

Division:

Result:



	Division	Computed	Hex	Binary
u	15213	15213	3B 6D	00111011 01101101
$u \gg 1$	7606.5	7606	1D B6	00011101 10110110
$u \gg 4$	950.8125	950	03 B6	00000011 10110110
$u \gg 8$	59.4257813	59	00 3B	00000000 00111011

Compiled unsigned division code

C Function

```
unsigned udiv8(unsigned x)
{
    return x/8;
}
```

Compiled Arithmetic Operations

```
shr1 $3, %eax
```

Explanation


```
# Logical shift
return x >> 3;
```

- Uses logical shift for unsigned
- For Java users: logical shift written as >>>

Signed power-of-2 divide w/ shift

- Quotient of Signed by Power of 2

- $s \gg k$ gives $\lfloor s / 2^k \rfloor$
- Uses arithmetic shift Operands:




Signed power-of-2 divide w/ shift

- Quotient of Signed by Power of 2

- $s \gg k$ gives $\lfloor s / 2^k \rfloor$
- Uses arithmetic shift Operands:
- Rounds wrong direction when $s < 0$.**

Division:

Result: $\text{RoundDown}(s/2^k)$




	Division	Computed	Hex	Binary
s	-15213	-15213	C4 93	11000100 10010011
s >> 1	-7606.5	-7607	E2 49	1 1100010 01001001
s >> 4	-950.8125	-951	FC 49	1111 1100 01001001
s >> 8	-59.4257813	-60	FF C4	11111111 11000100

Correct power-of-2 divide

- Quotient of negative number by power of 2
 - We want $\lceil s/2^k \rceil$ (round toward 0)
 - We compute it as $\lfloor (s + 2^k - 1)/2^k \rfloor$
 - In C: $(s + (1 \ll k) - 1) \gg k$
 - Biases the dividend toward 0

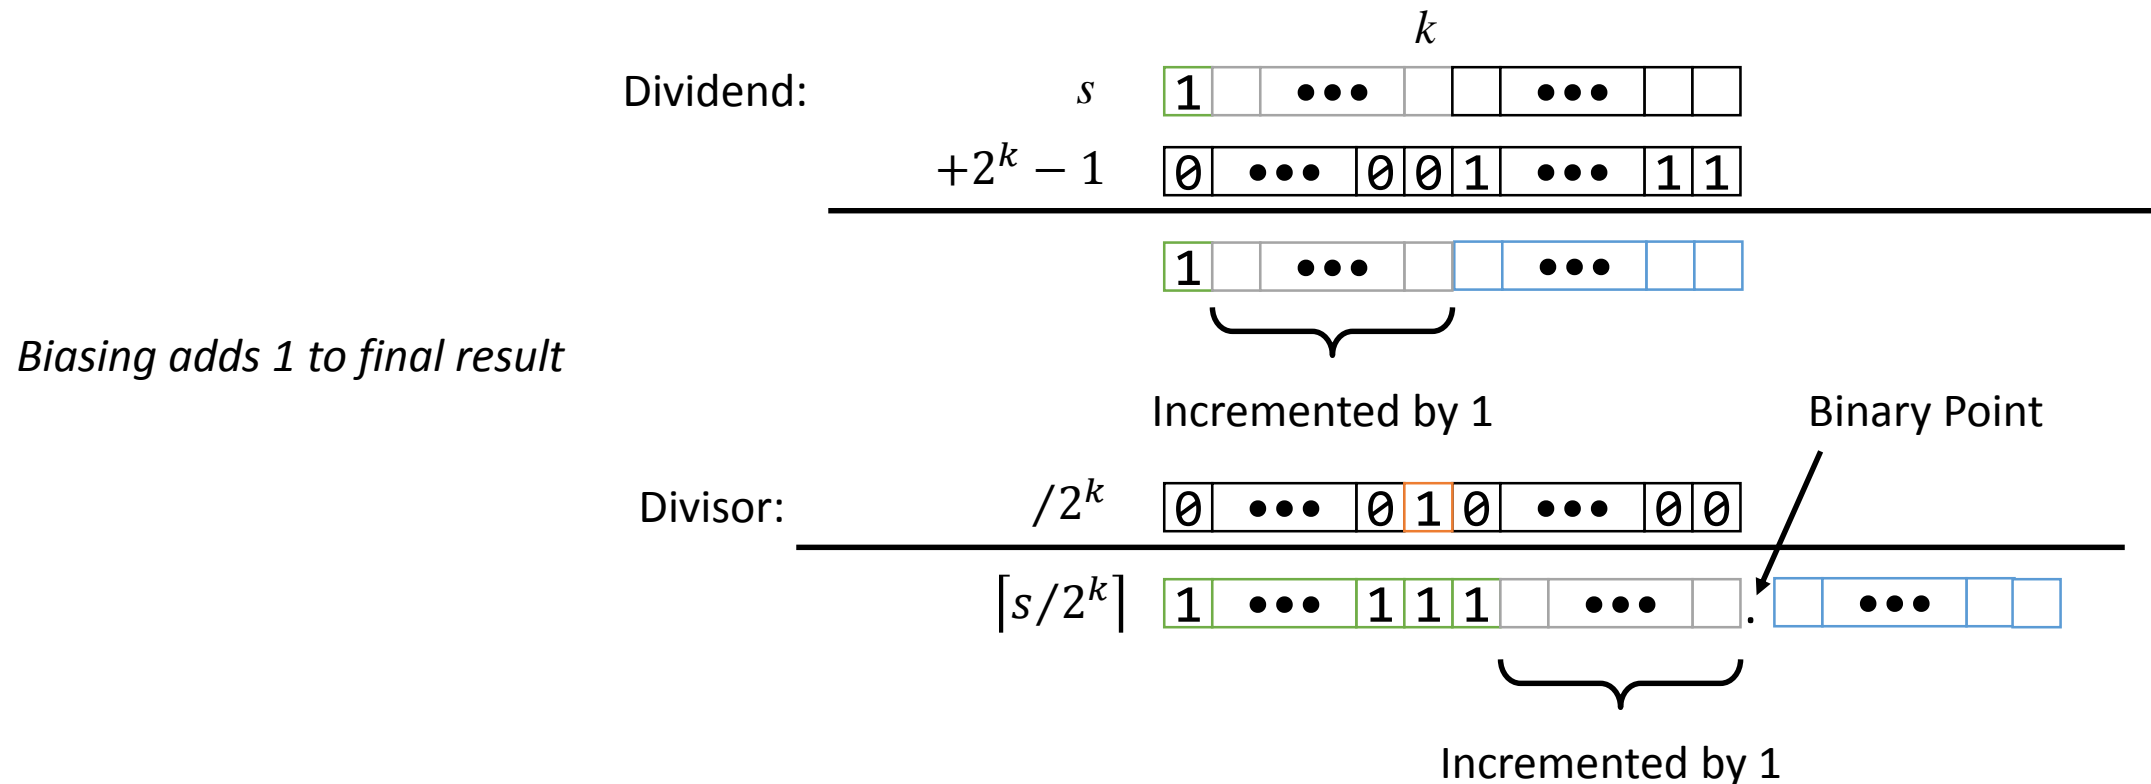
• Case 1: No rounding

Biasing has no effect



Correct power-of-2 divide (Cont.)

- Case 2: Rounding:



Compiled signed division code

- Uses arithmetic shift for int
- For Java users
 - Arith. shift written as >>

Explanation

```
if x < 0
    x += 7;
# Arithmetic shift
return x >> 3;
```

C function

```
int idiv8(int x)
{
    return x/8;
}
```

Compiled arithmetic operations

```
testl %eax, %eax
js    L4
L3:
    sarl $3, %eax
    ret
L4:
    addl $7, %eax
    jmp  L3
```


Summary

- Signed/unsigned multiply:

$$x * 2^k = x \ll k$$

- Unsigned divide:


$$u / 2^k = u \gg k$$



- Signed divide:

$$s / 2^k = s \gg k \quad \text{for } s > 0$$

$$s / 2^k = s + (2^k - 1) \gg k \quad \text{for } s < 0$$



3.6: C Integer puzzles

Computer Architecture and Systems Programming

Integer C Puzzles

- Assume 32-bit word size, two's complement integers
- Implementation defined by hardware
- For each of the following C expressions, either:
 - Argue that it is true for all argument values
 - Give an example where it is not true

Initialization:

```
int x = foo();  
int y = bar();  
unsigned ux = x;  
unsigned uy = y;
```

- $x < 0 \Rightarrow ((x*2) < 0)$
- $ux \geq 0$
- $x \& 7 == 7 \Rightarrow (x \ll 30) < 0$
- $ux > -1$
- $x > y \Rightarrow -x < -y$
- $x * x \geq 0$
- $x > 0 \&\& y > 0 \Rightarrow x + y > 0$
- $x \geq 0 \Rightarrow -x \leq 0$
- $x \leq 0 \Rightarrow -x \geq 0$
- $(x \mid -x) \gg 31 == -1$
- $ux \gg 3 == ux/8$
- $x \gg 3 == x/8$
- $x \& (x-1) != 0$